



# **Tag You're It!**

## **Revisiting the Reality of DNSSEC Keytags**

Roland van Rijswijk-Deij

# Introduction

- **DNSSEC validation** requires a resolver to **match signatures to keys**  
→ it is **inefficient to check** a signature **against every DNSKEY** in a zone
- To enable **fast matching**, DNSSEC has the notion of **keytags**  
(first introduced in draft 01 of what became RFC 2535 in August 1997)
- These are **16-bit values** included in RRSIG records and can **help** a resolver  
**find a matching key**
- They are **only a hint**; the **idea is** that it is **unlikely for keytag collisions to occur** in a DNSKEY set, so they **help** match keys to signatures **in most cases**

# Introduction

- **Some years ago**, Roy Arends presented at **DNS OARC** on the **curious case of the unused keytags** [1]
- He generated **large numbers of keys**, and found that **only between a quarter and half out of the 65536 possible keytags occurred**
- **Due to mathematical properties of RSA keys and** how the **keytag algorithm** computes the tag from the key's RDATA

[1] <https://slideplayer.com/slide/10329289/>

# In this talk

- While - with a generous community effort - Roy managed to explain **why certain keytags do not occur in theory**
- In this talk we look at **what keytags occur in practice**
- And **draw lessons** from this **for protocol design** (or: why we should have picked a different, simple algorithm for keytag computation)

# Quick refresher: keytag algorithm

- Assuming "self" is a Python DNSKEY object, this is the algorithm:

```
def keytag(self):
    acc = int(0)

    wire = self.towire()

    for i in range(0, len(wire)):
        if i & 1 == 1:
            acc += wire[i]
        else:
            acc += wire[i] << 8

    acc += (acc >> 16) & 0xffff

    return acc & 0xffff
```

- Basically accumulate even bytes in the lower 8 bits, odd bytes in the high 8 bits, and do some fiddling with carry bits
- The outcome of this algorithm highly depends on the amount of structure and predictability in the input!

# RSA keys have a lot of structure

- **Skipping** the **details** (which can be found in Roy Arends' OARC presentation), **RSA keys have a lot of structure**
- E.g. due to the **modulus always** being **odd**, **use of safe primes** in key generation, strong **preference for** certain **public exponents**, ...
- Also: **flags, protocol version and algorithm** are **also included** in the computation!
- In the **OARC presentation**, there was talk of **either 16384 or 32768 possible keytags** (but **did not take every case into consideration**)

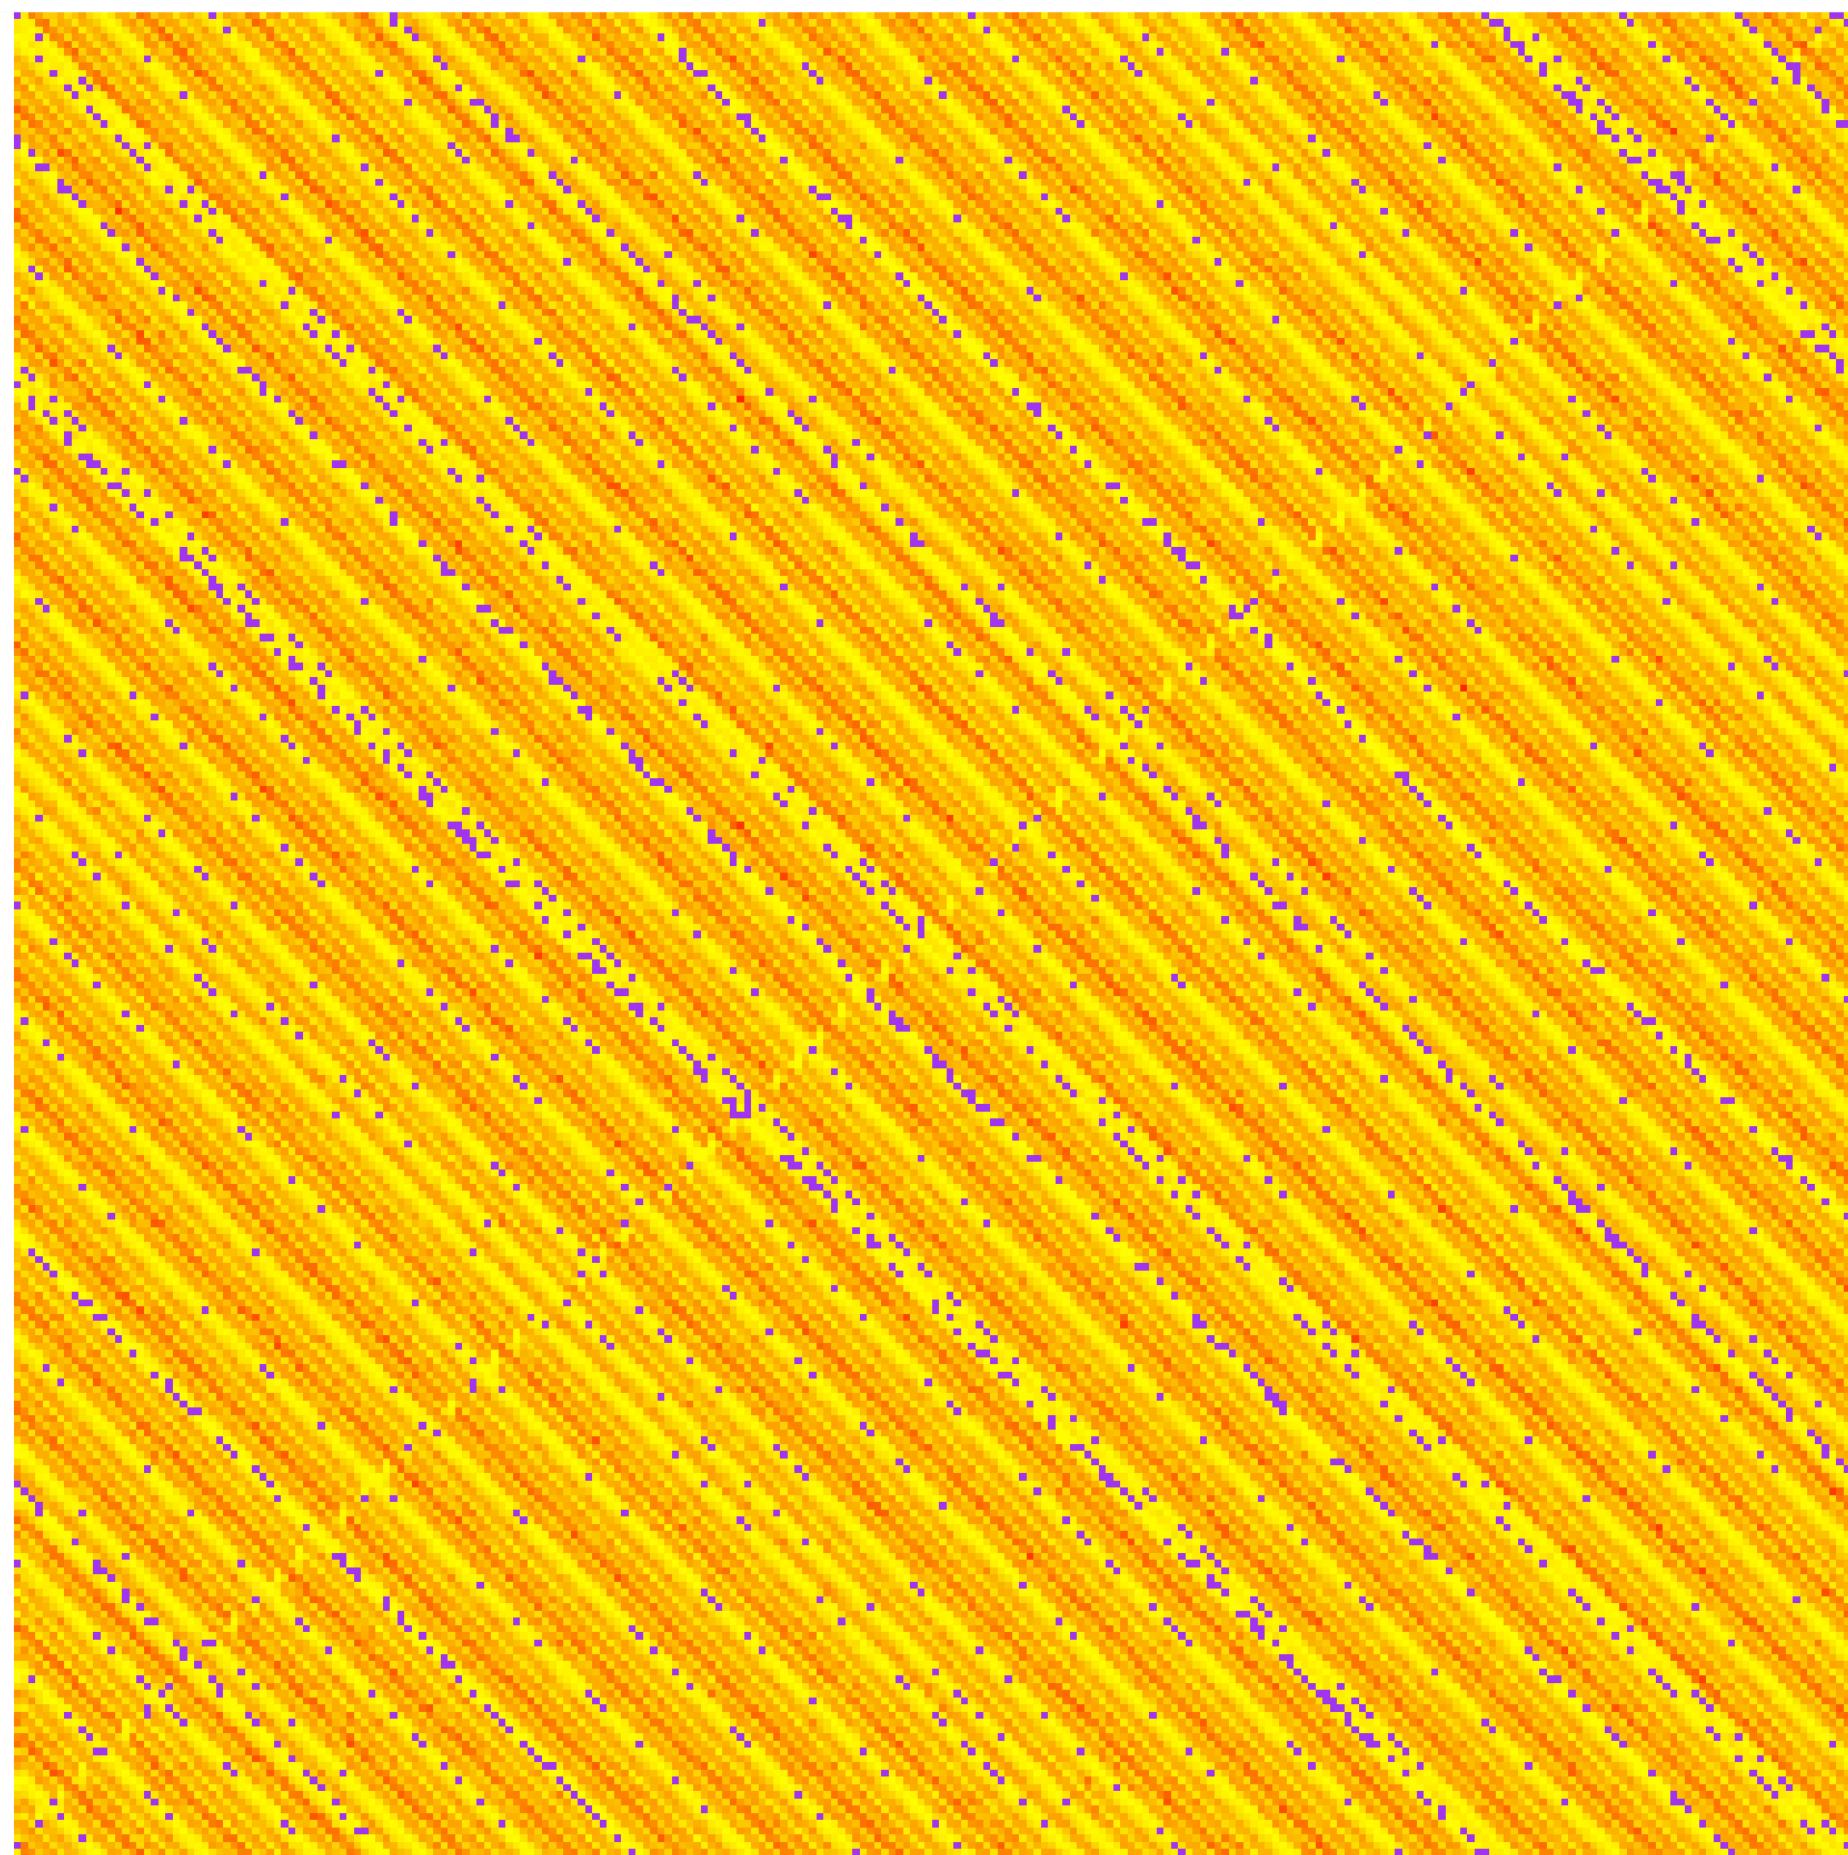
# What happens in practice?

- **Experiment** discussed at OARC relied on generating lots of **keys**
- What we wanted to know: **what happens in the wild?**
- We took one recent day of **data from OpenINTEL for .com and .nl** and looked at two things:
  - **Keytags in the wild** for RSA and ECDSA
  - **Keytag collisions in the wild** for RSA and ECDSA

# What we expect to find

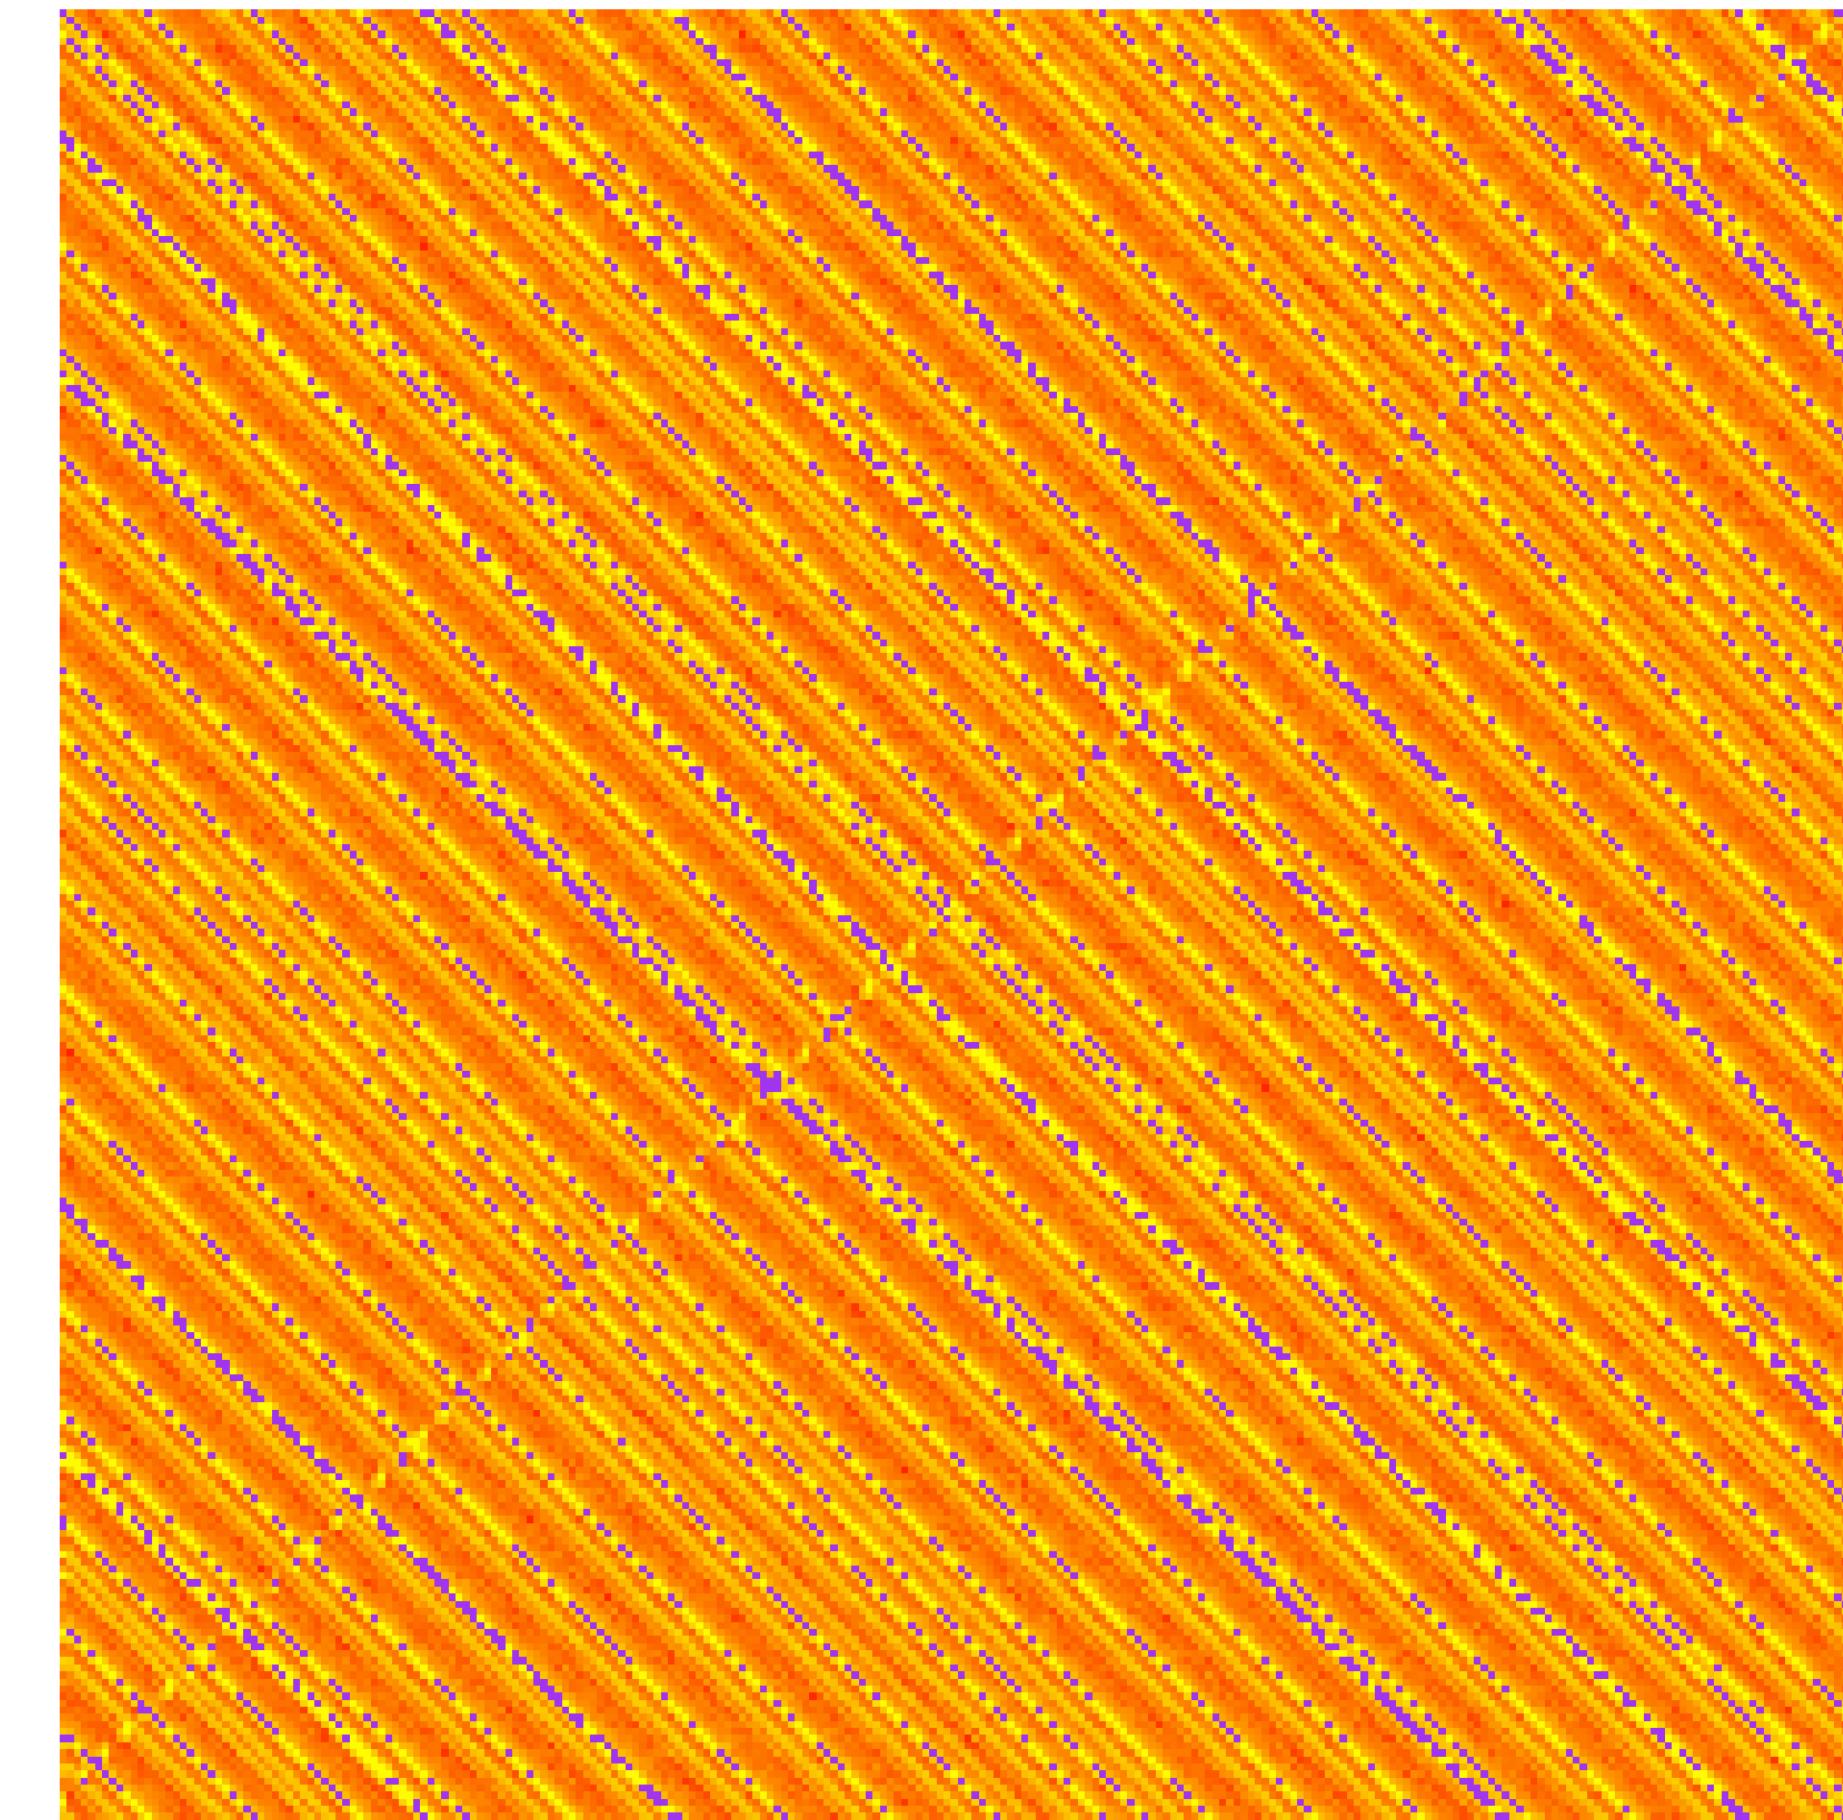
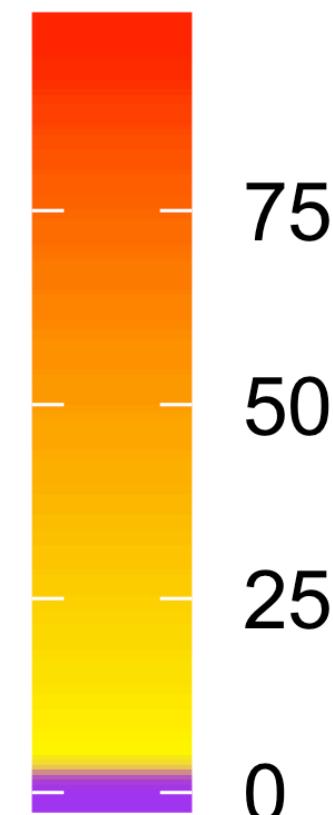
- As stated, RSA keys contain a lot of structure, and the keytag algorithm output is influenced by structure
- We **expect to find:**
  - That **certain keytags** are **much more common for RSA** keys
  - That **keytag collisions occur in the wild for RSA** keys
  - That the **occurrence of keytags for ECDSA** is much more **uniformly distributed** due to a lack of structure in ECDSA keys

# Heatmaps for RSA keytags



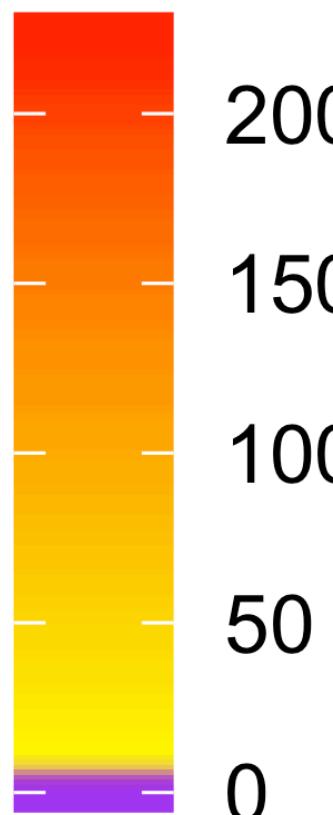
.com

value

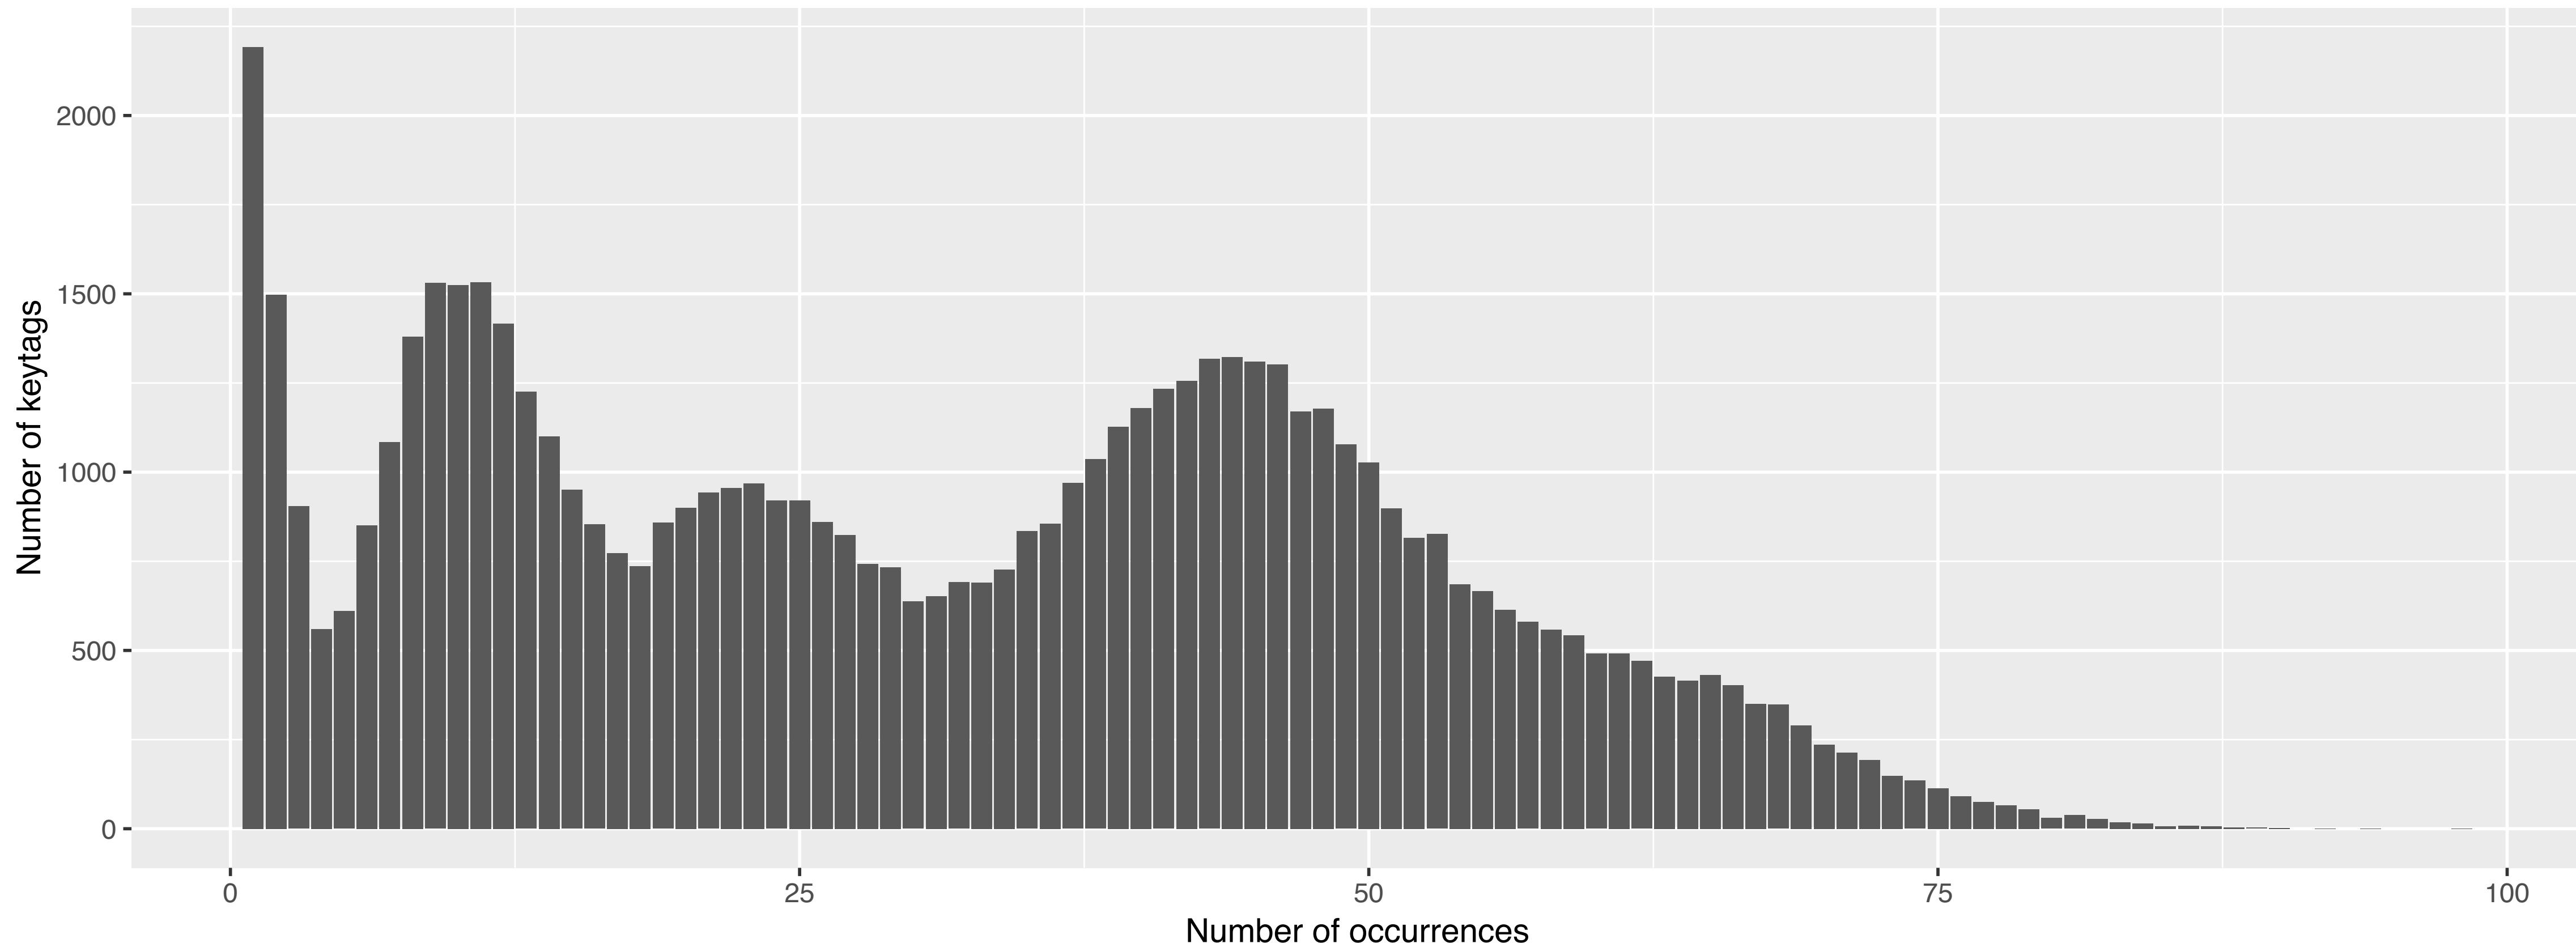


.nl

value

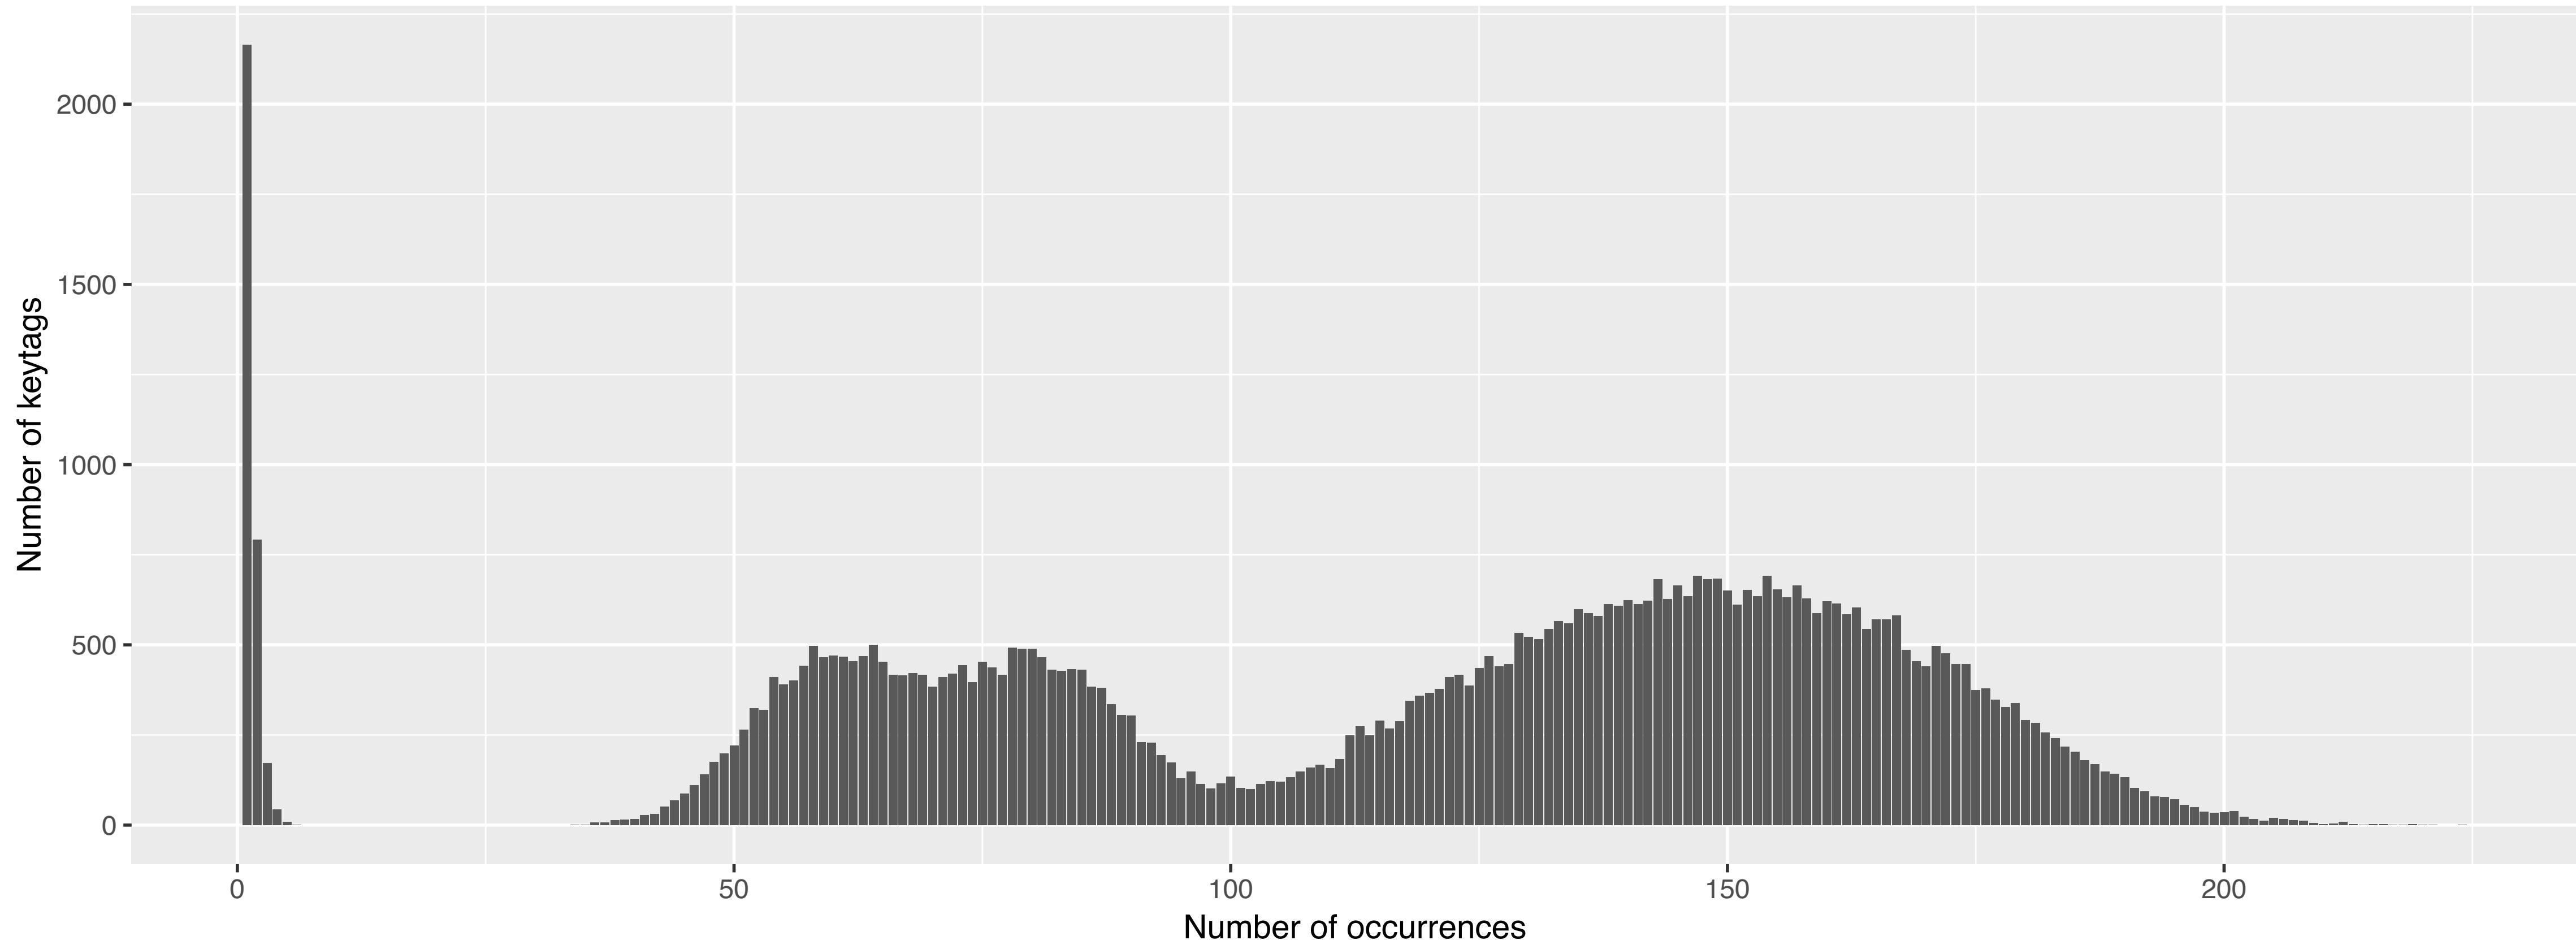


# Distribution of RSA keytags in .com



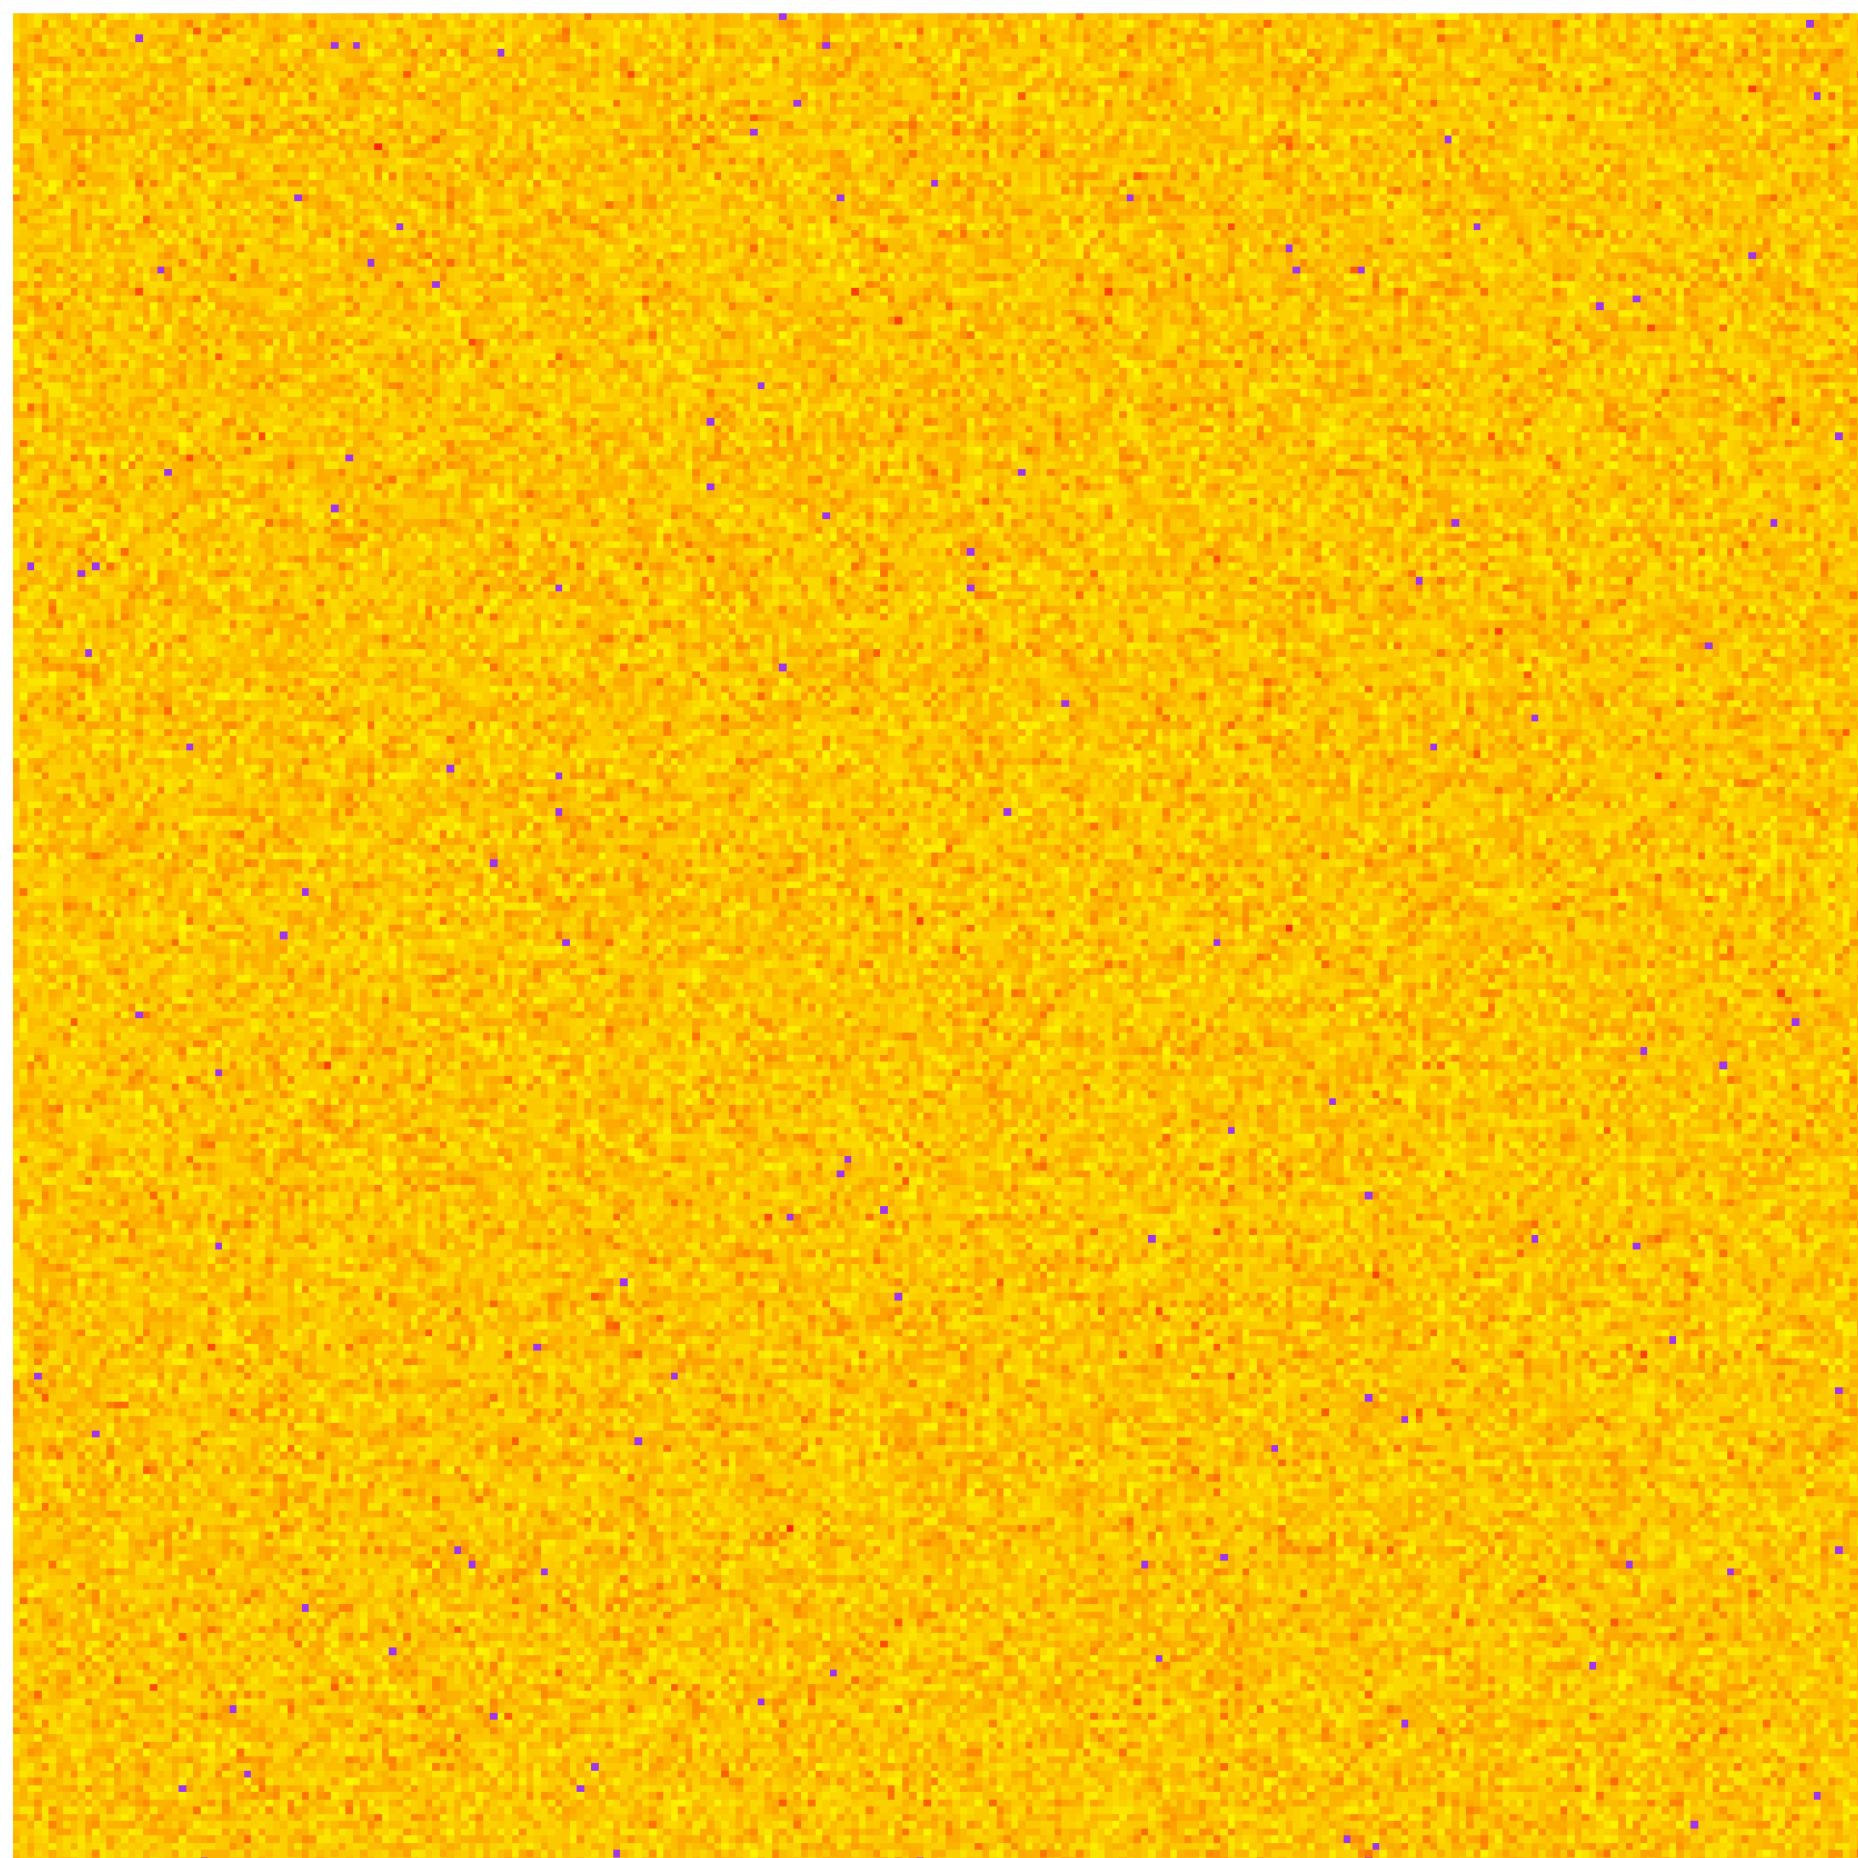
- Takeaway: some keytags occur only once, a few keytags occur over 75 times more often than that

# Distribution of RSA keytags in .nl

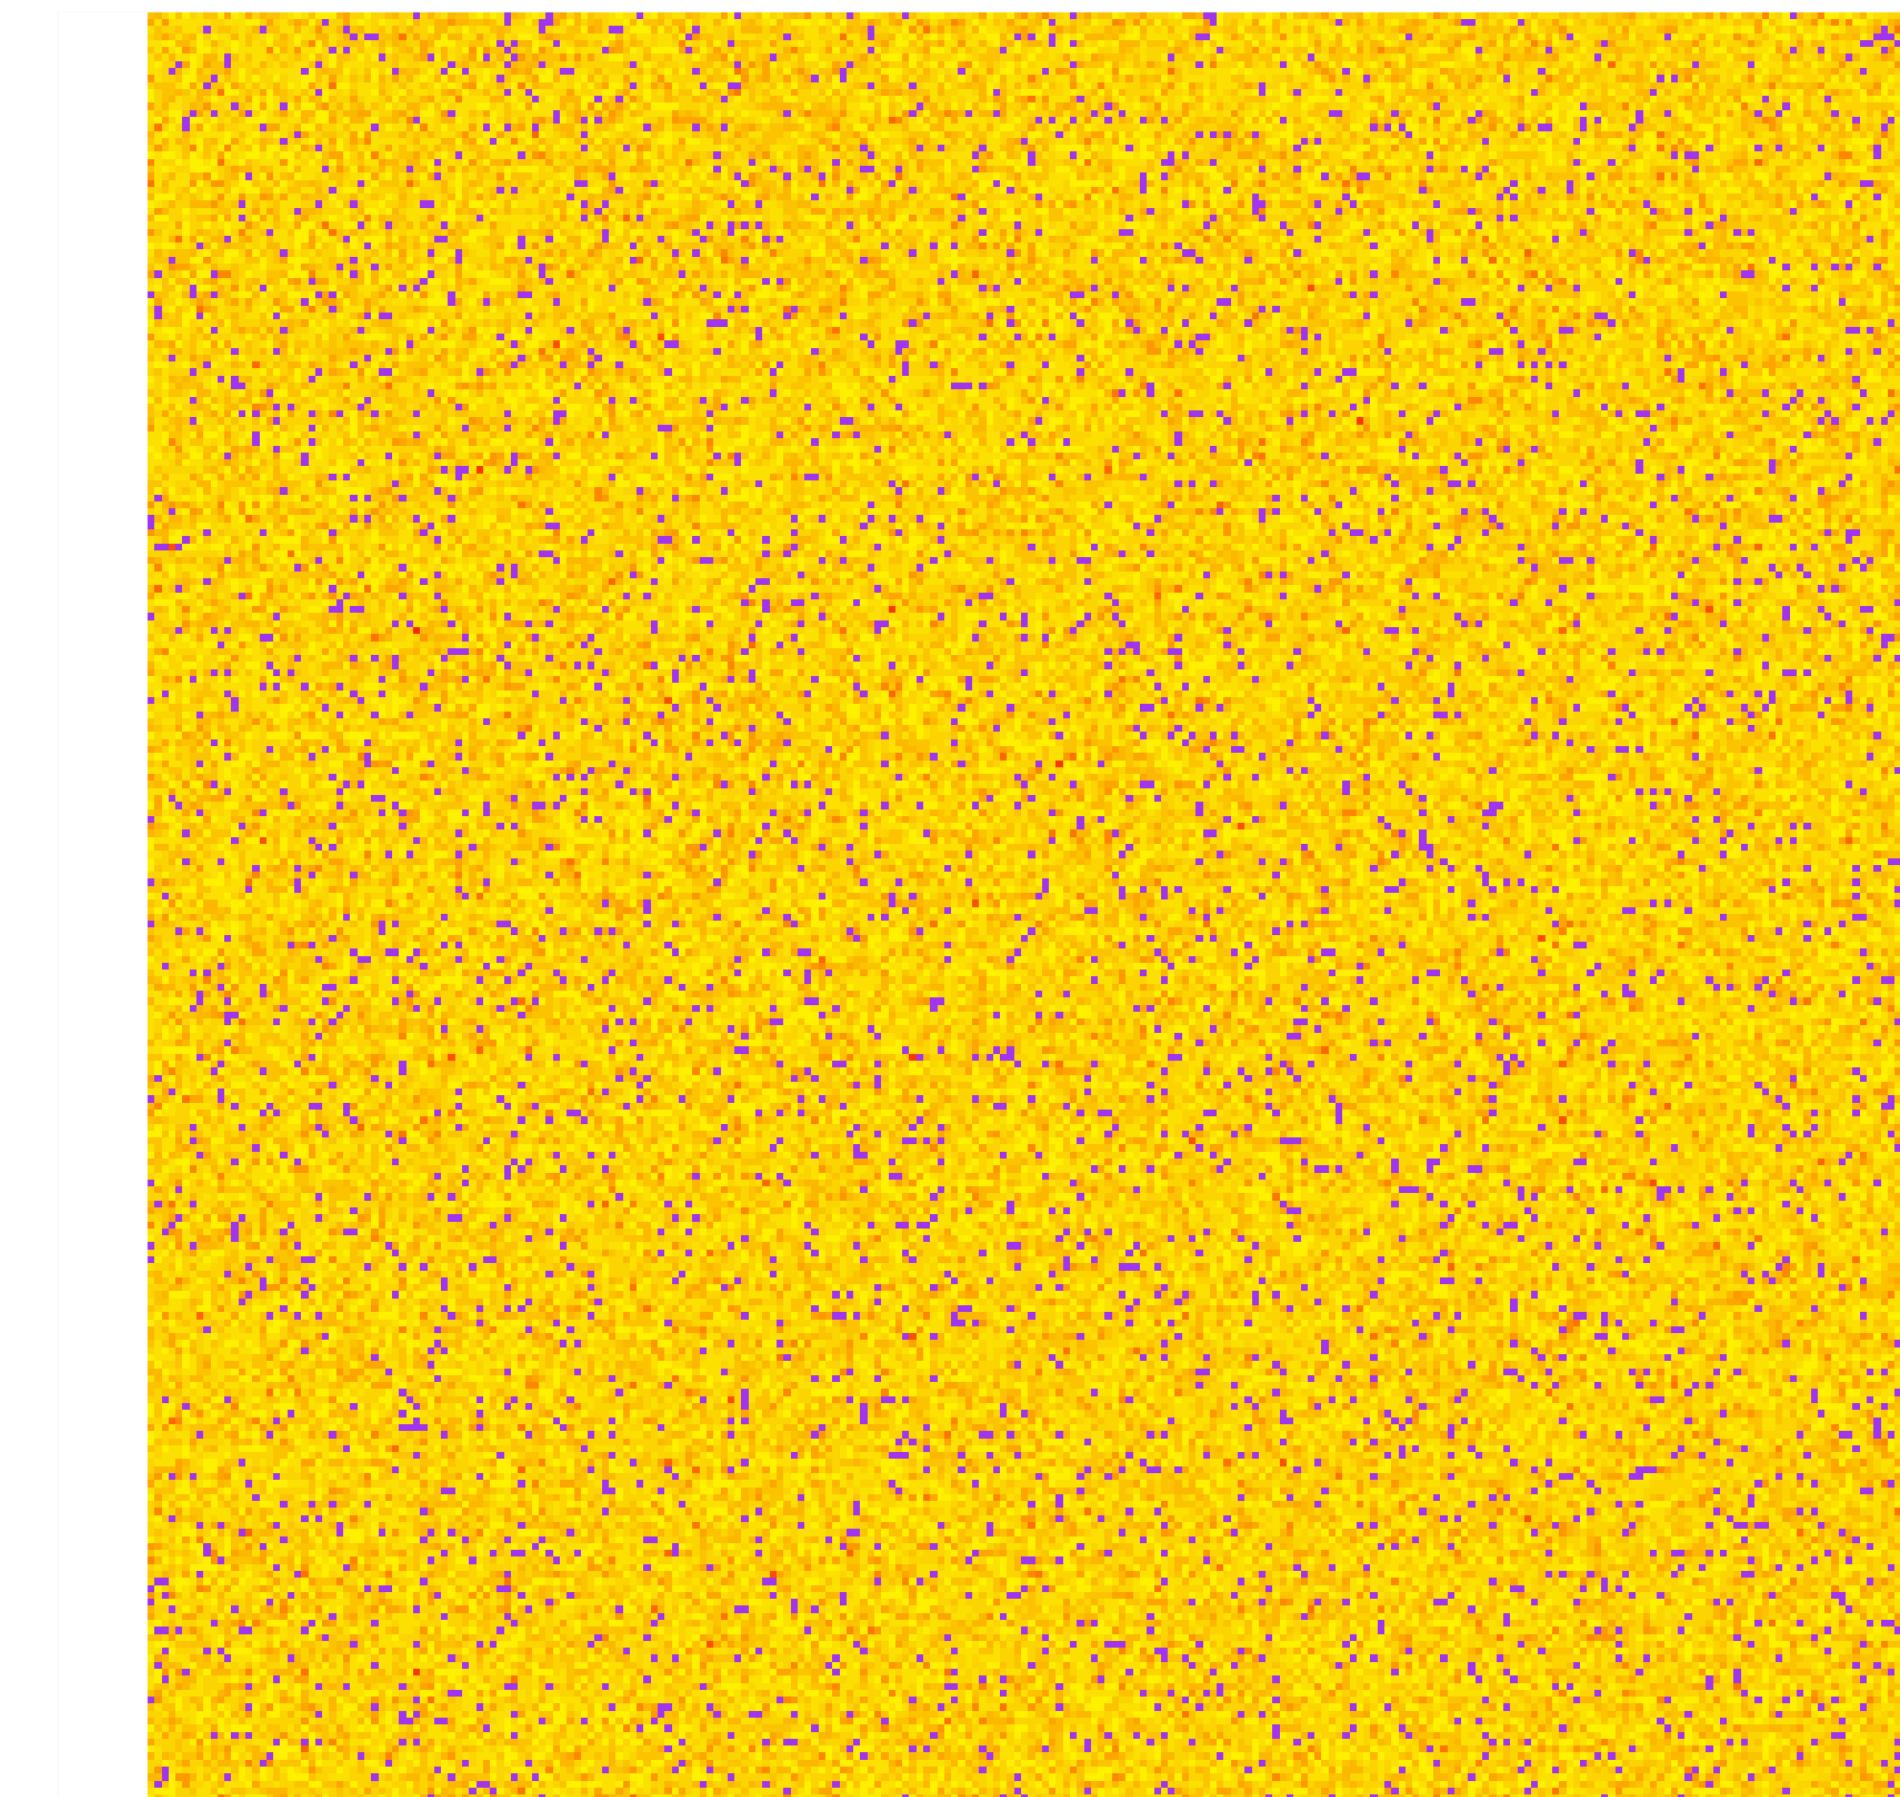


- Why the difference? Different distribution in algorithms and key sizes!

# So what about ECDSA?

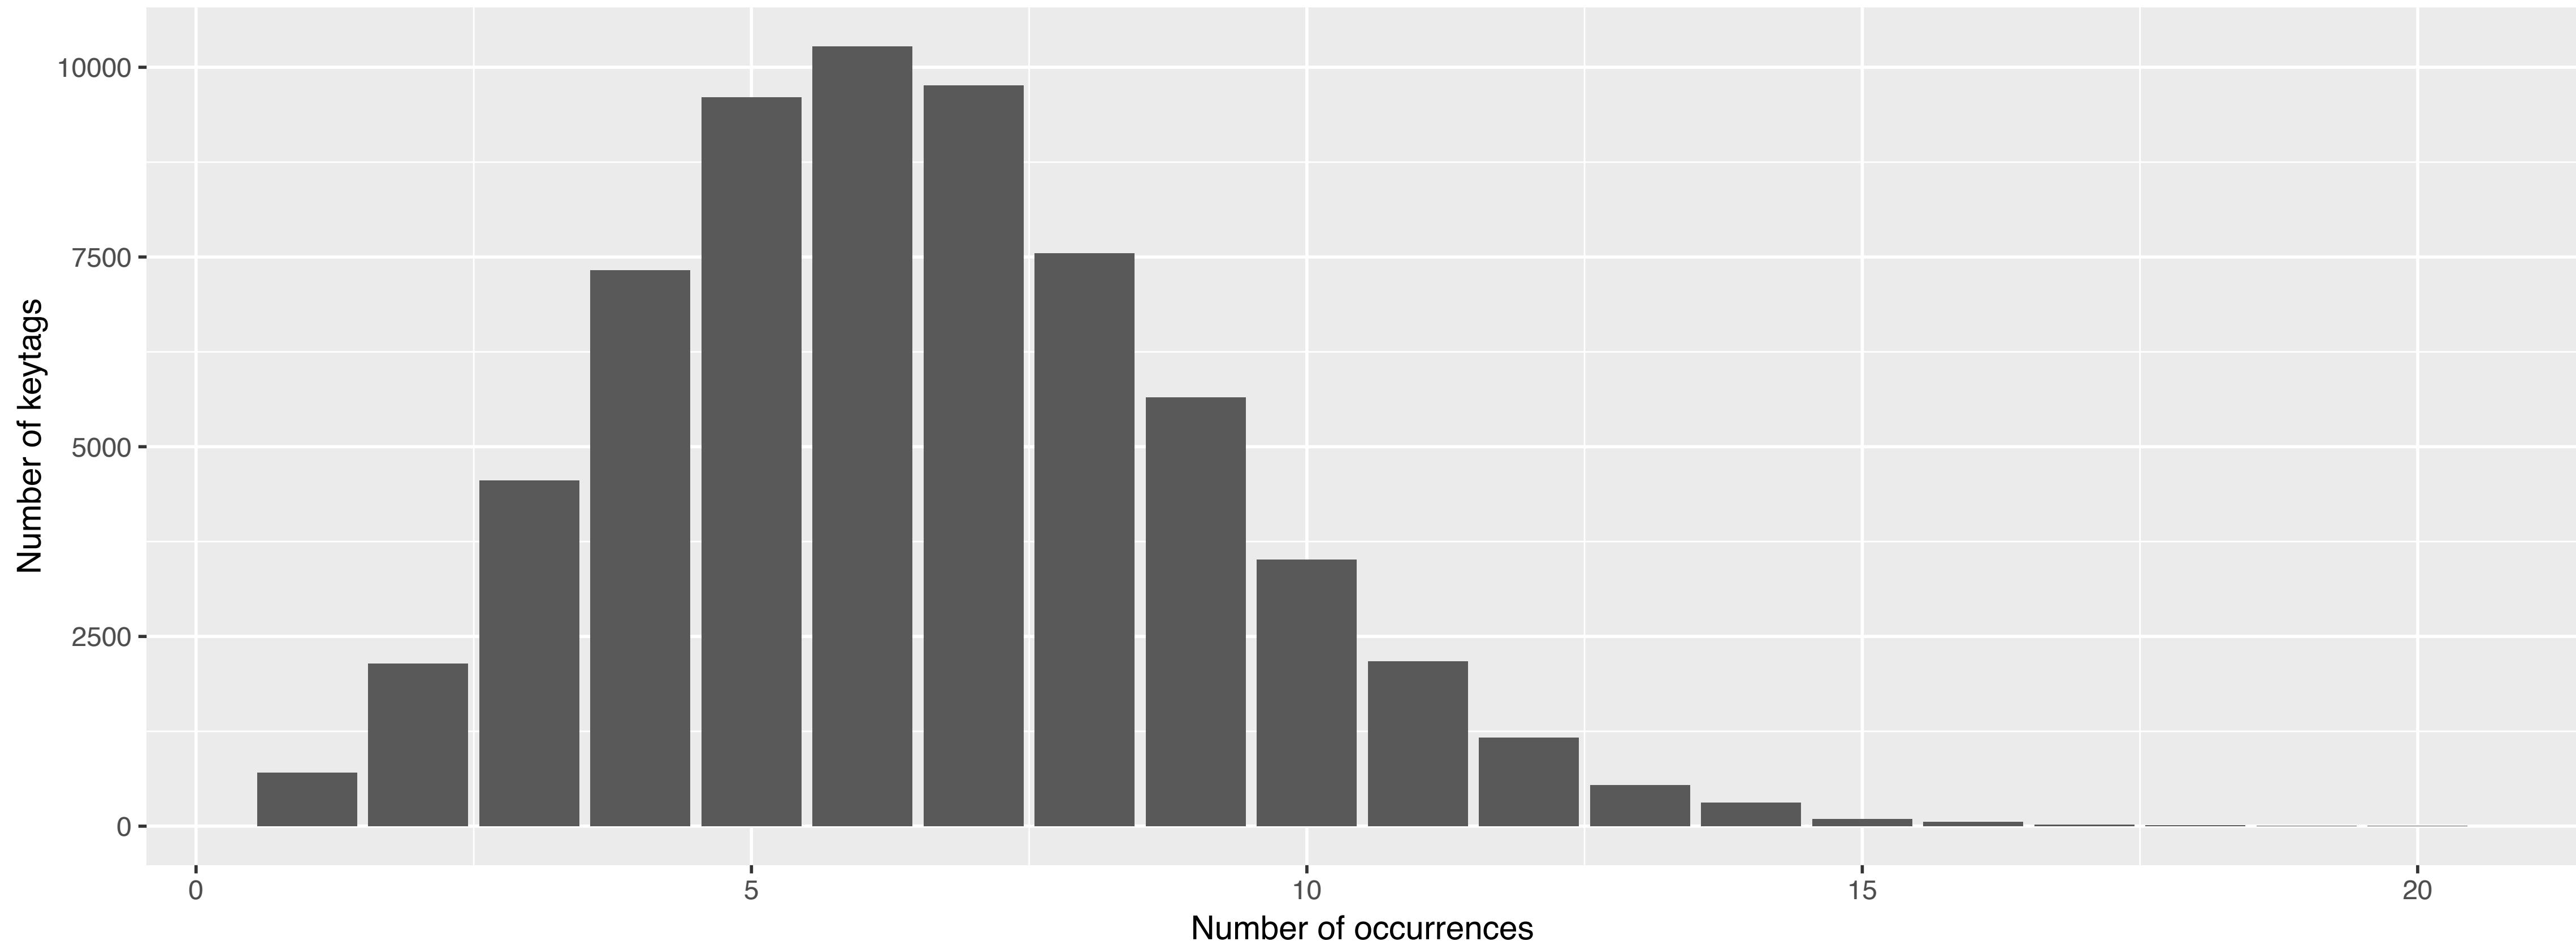


.com



.nl

# ECDSA keytag distribution

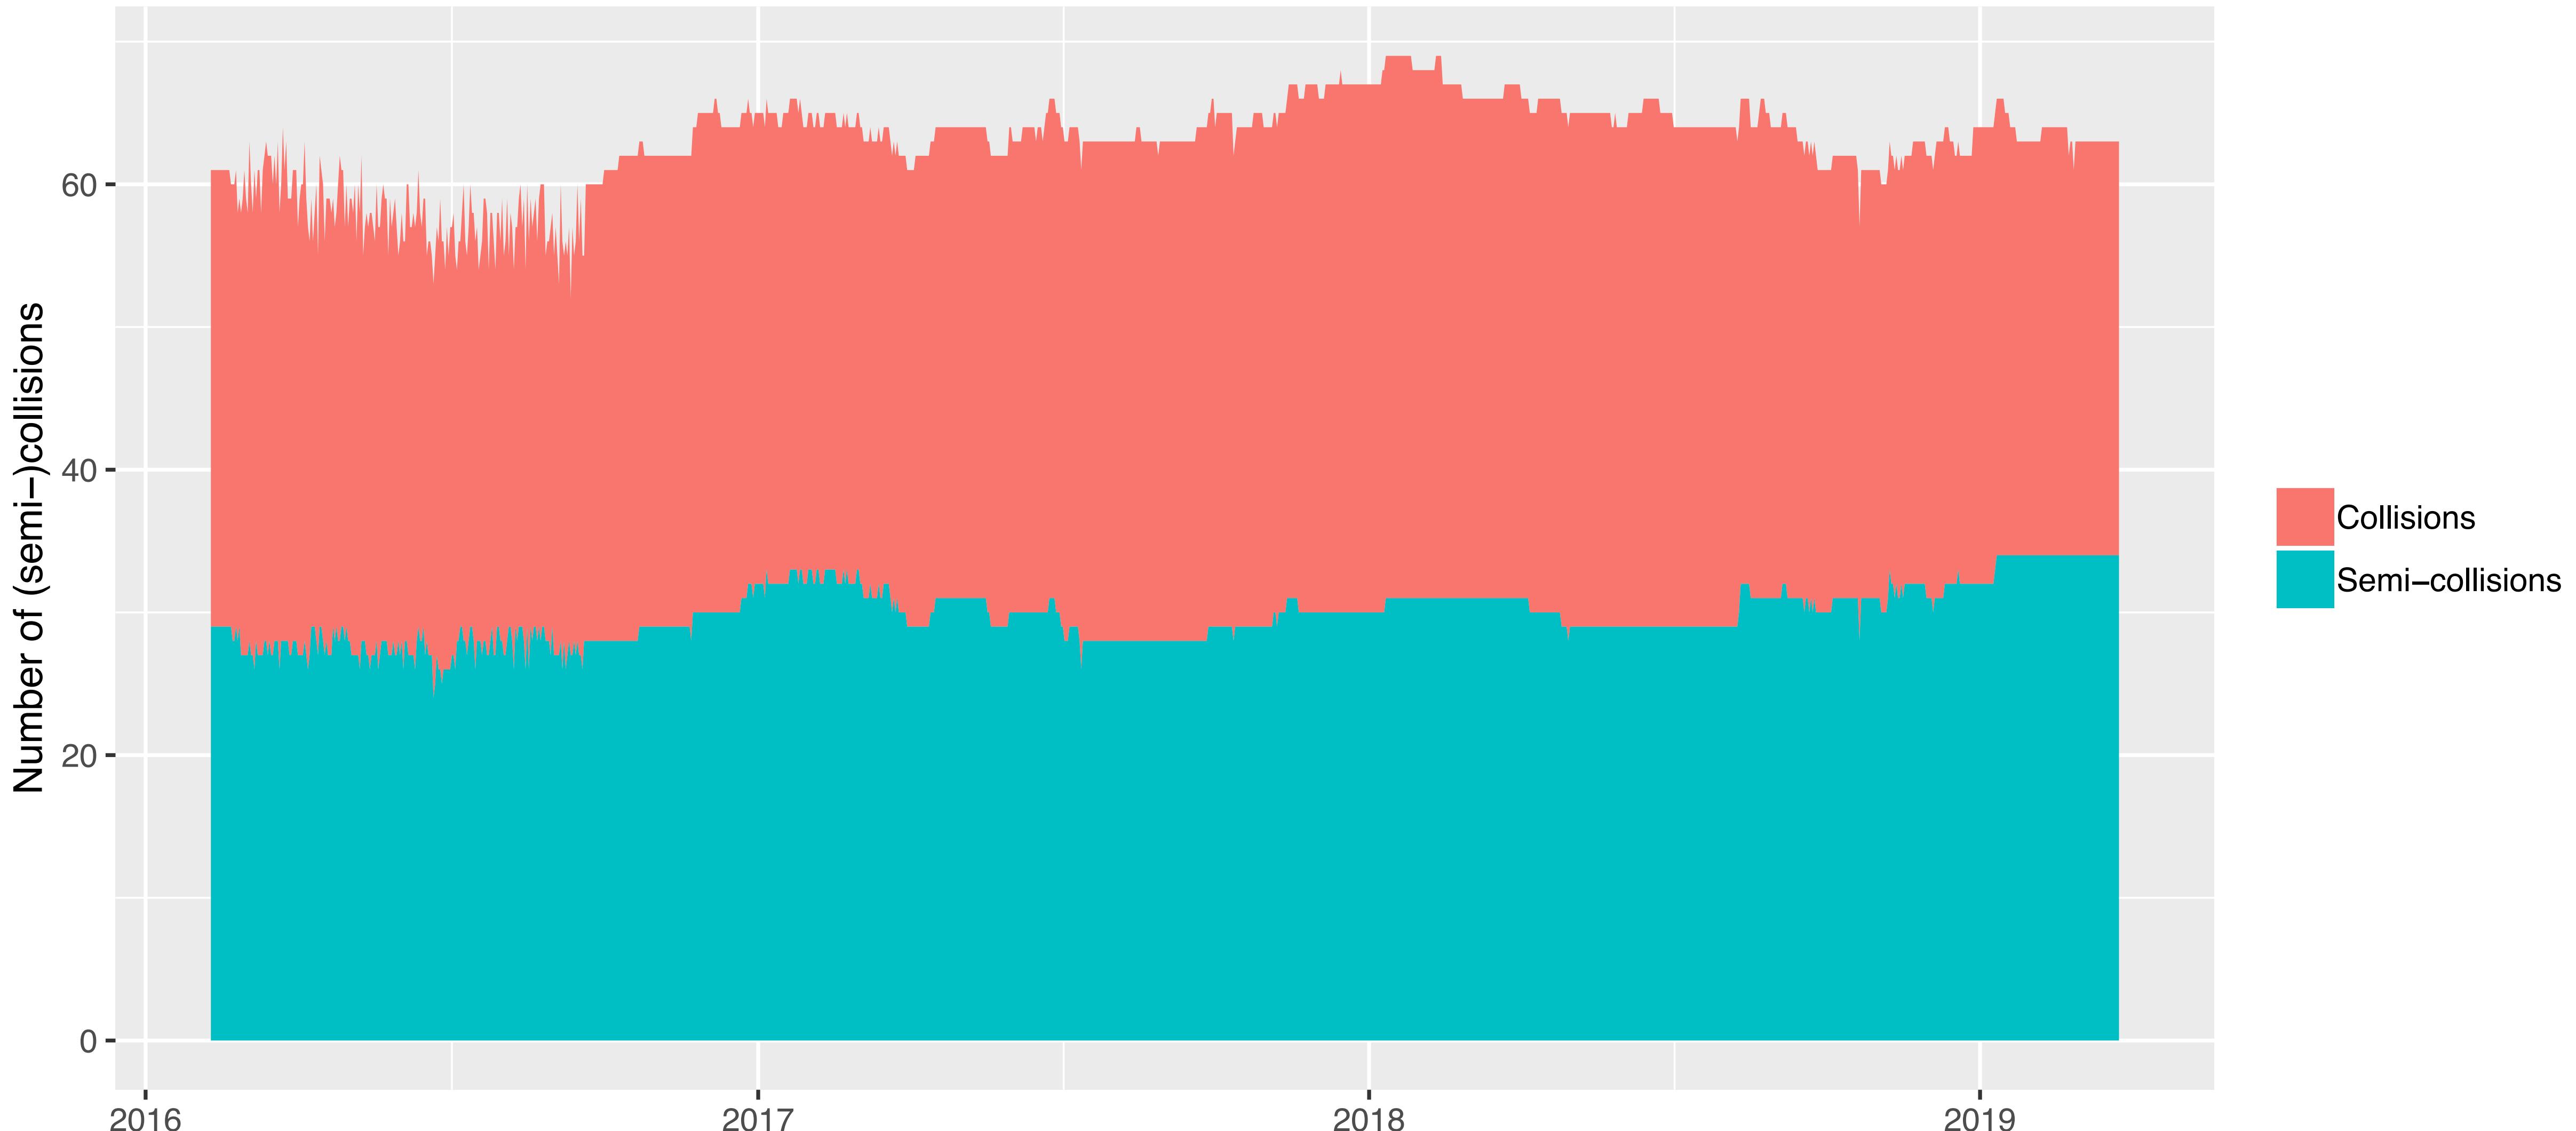


- This **looks much more like a Gaussian** distribution (but isn't quite)

# Collisions

- **Searched for collisions in** OpenINTEL, by computing keytags for DNSKEY RRsets for all signed domains in **.nl** over **3 years of data**
- Distinguish between **two types of collisions:**
  - "**Real**" collisions, where two or more **different keys** in the keyset of the **same algorithm and size have the same keytag**
  - "**Semi**" collisions, where two or more different keys in the keyset have the **same key tag, but** have a **different size or algorithm**

# Collisions over time



# How rare are collisions?

- If keytags were uniformly distributed, we would expect the birthday paradox to apply, with  $n$  the number of keys and  $d$  the number of possible keytags; table shows this theoretical vs. actual probability

#keys in DNSKEY set	Theoretical probability	Observed probability
2	0,00153%	0,00122%
3	0,00458%	0,00322%
4	0,00915%	0,01616%
<b>5 or more</b>	0,01526%	0,00148%

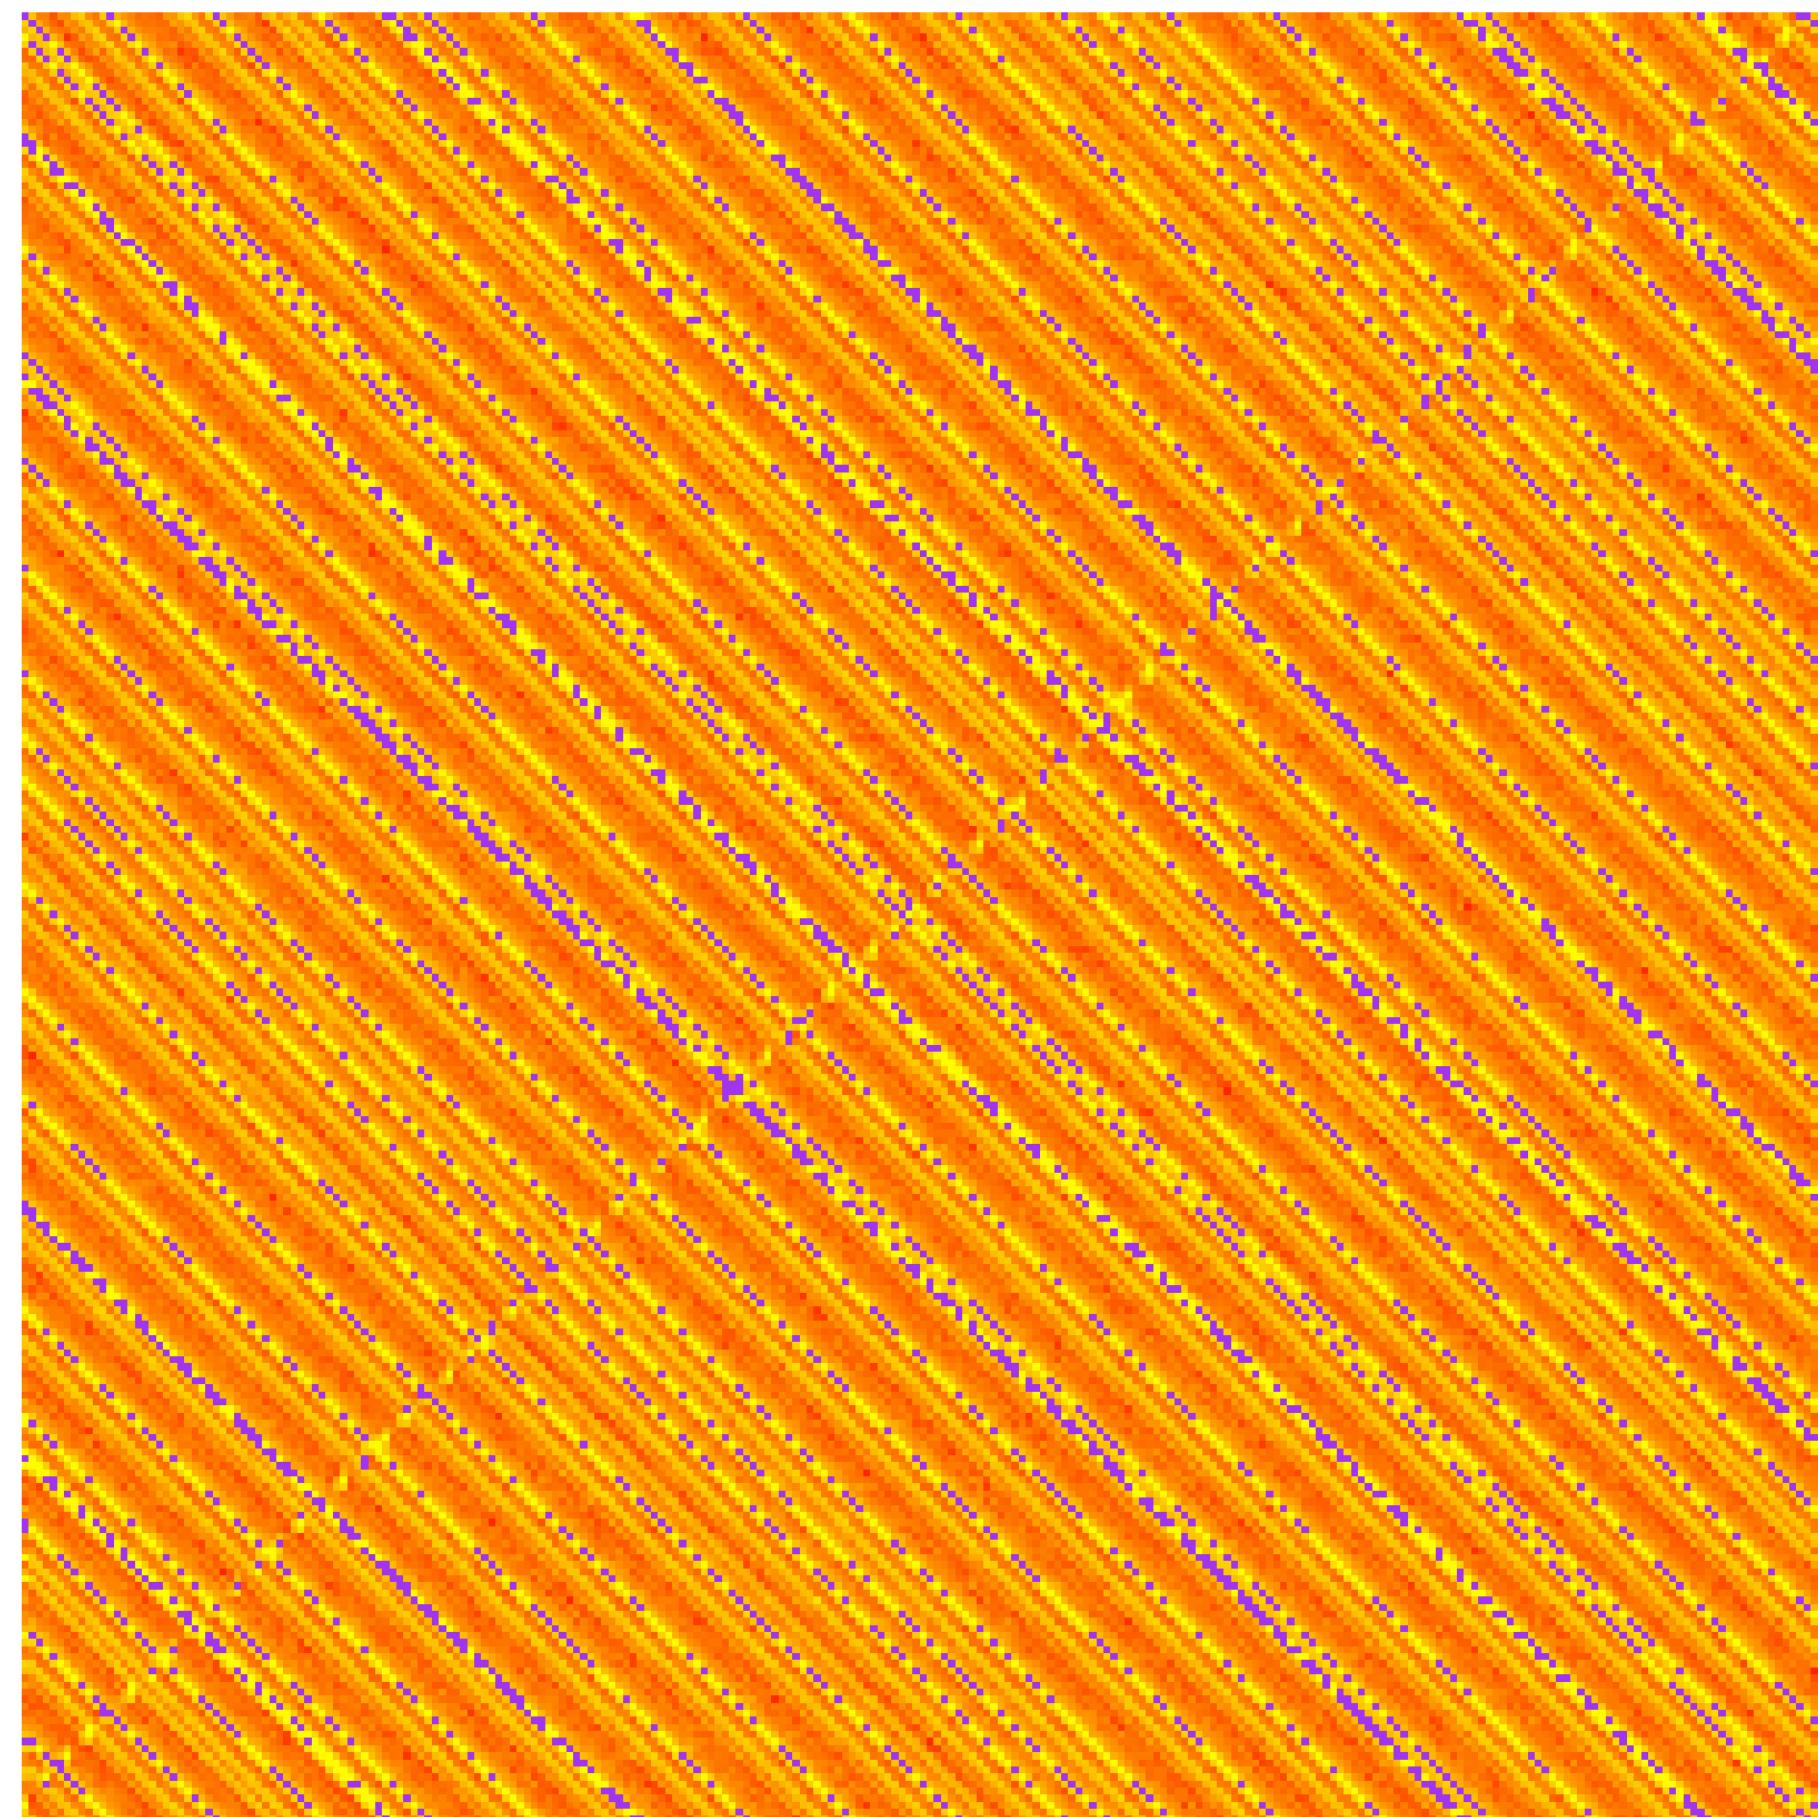
# Real-world impact

- **Keytags** are **only intended as a hint** for resolvers; they should **never solely** rely on them to **identify** the **correct key for** signature **validation**
- While very rare in practice, **collisions have a real impact**
- A **collision forces resolvers to validate signatures against all keys** that have a matching key tag until the correct key is found
- This **will lead to extra** CPU intensive **cryptographic operations**, that cause a small, but **quantifiable increase in load** on validating resolvers

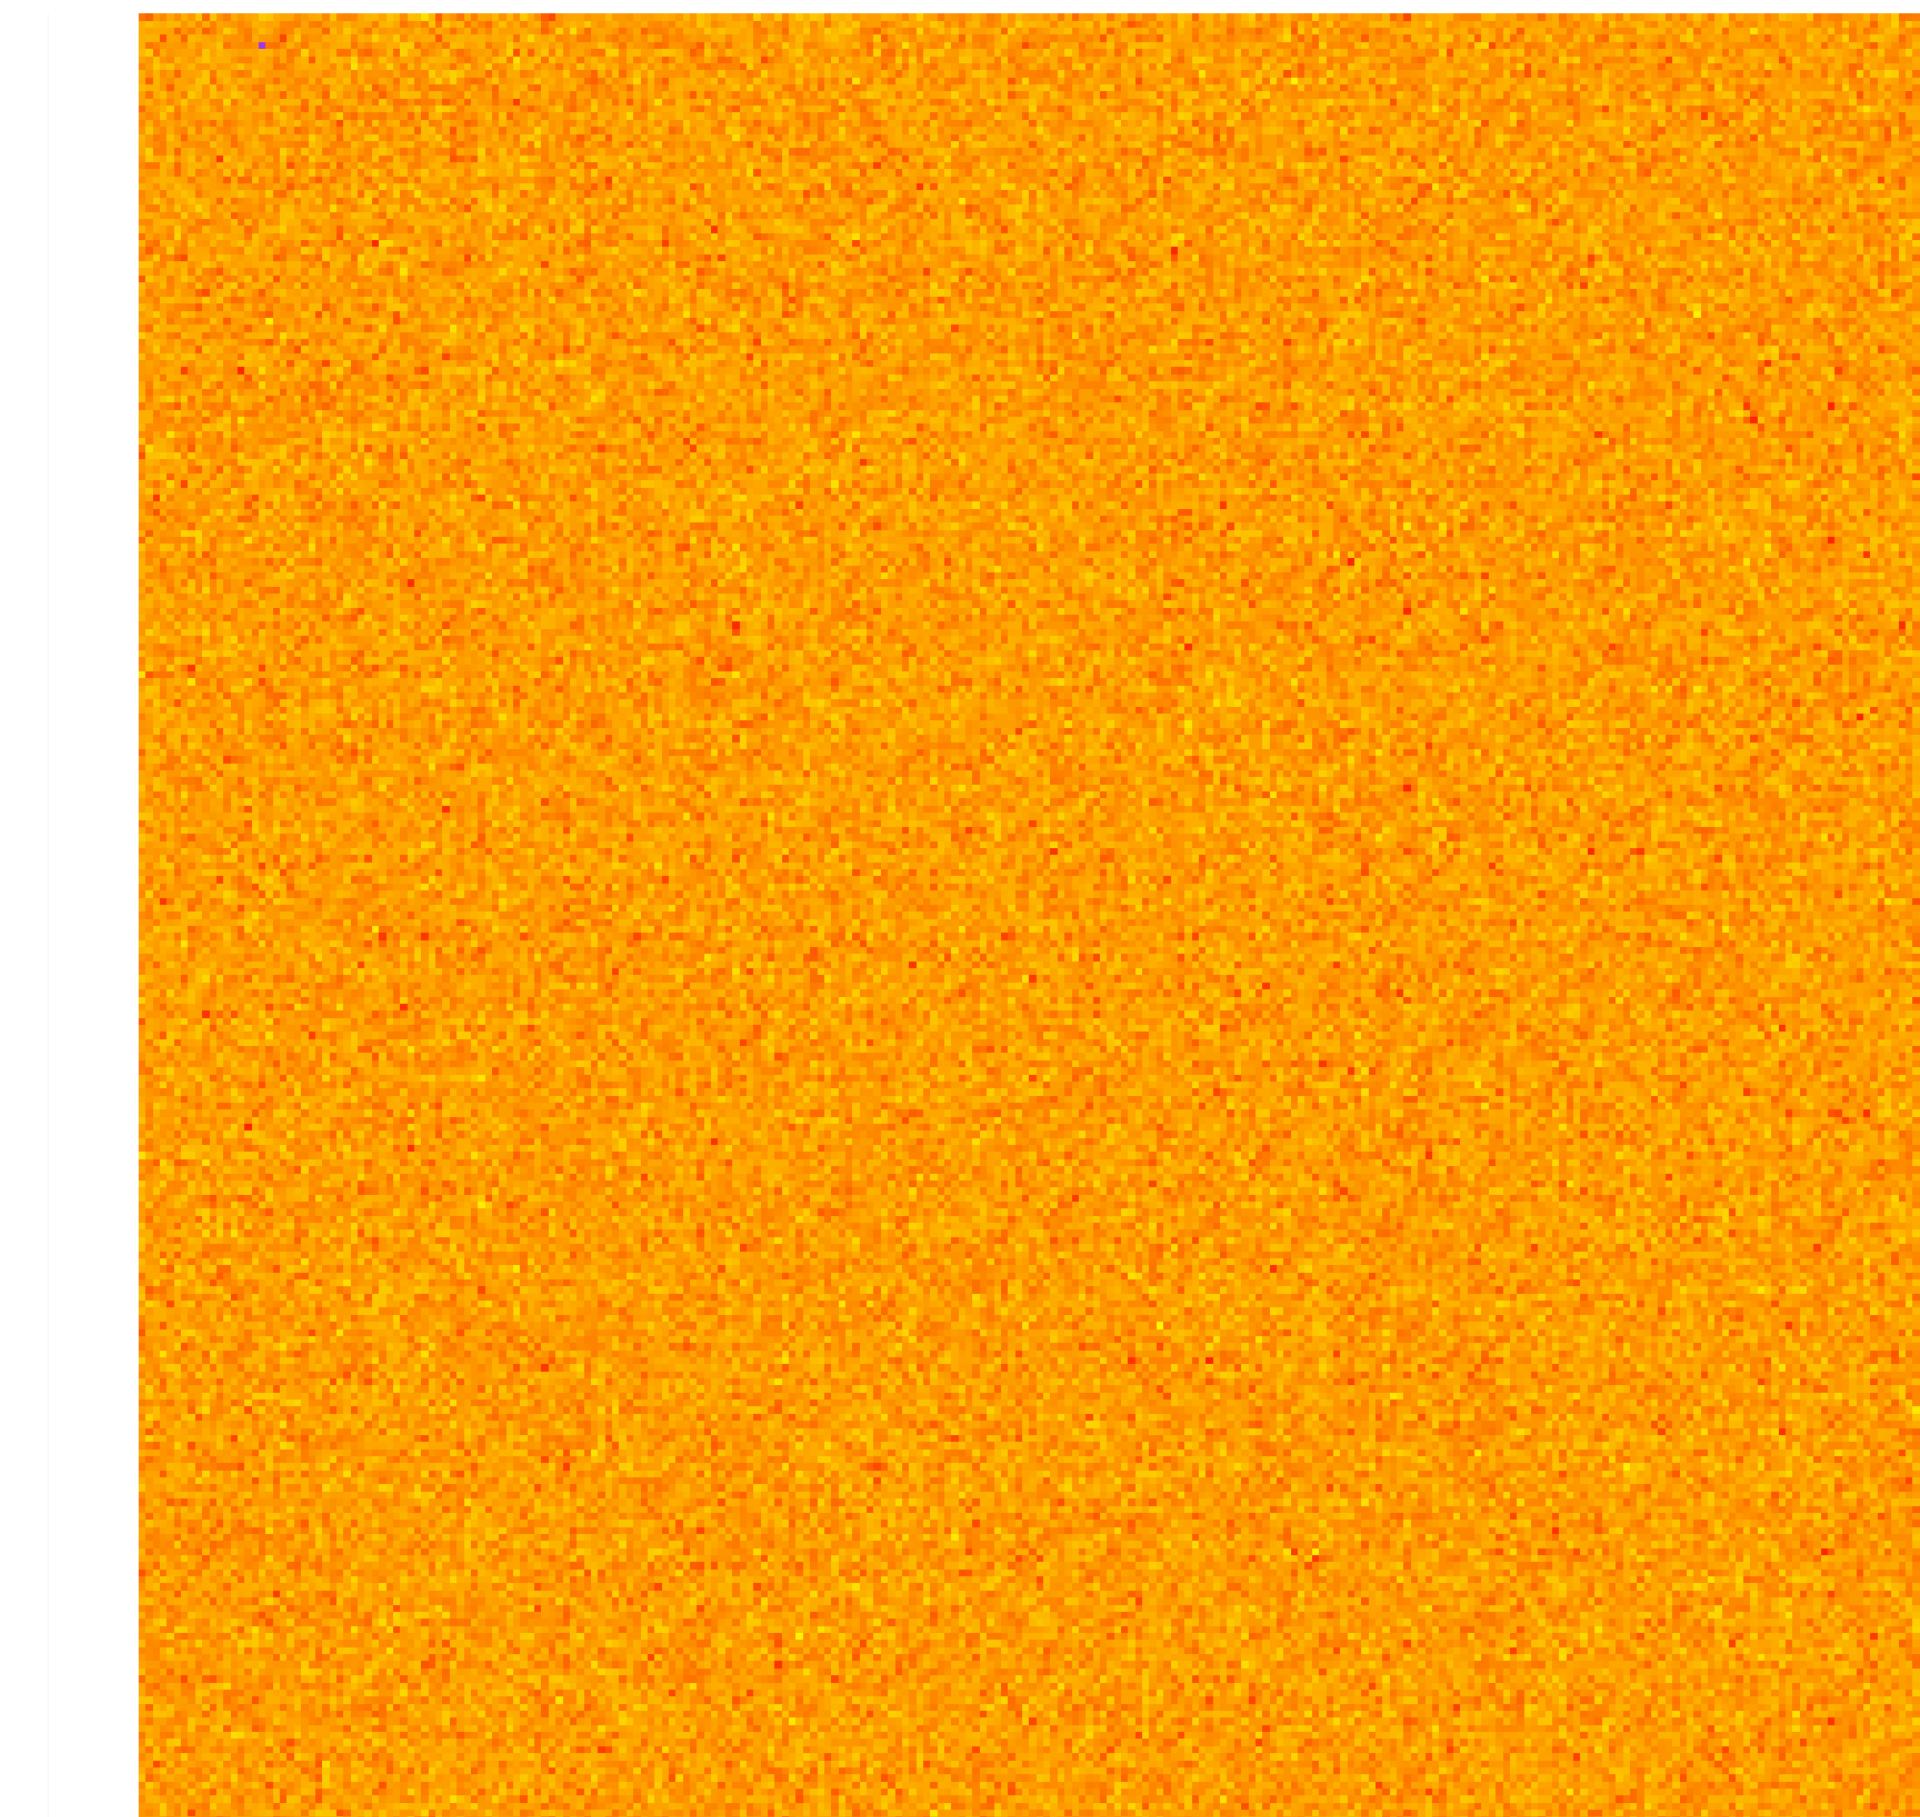
# What could we do differently?

- Clearly, the **keytag algorithm appears not optimal** for its purpose, **but** appears to make **collisions less likely than expected**
- So **what** would happen **if** we used something that produces a **random uniform output**, regardless of any structure in the input?
- (Cryptographic) **hash functions have this property, but** are likely **much more computationally intensive**
- Yet there is a **middle ground:** what if we used **CRC16?**

# Heatmap: keytag vs. CRC16 for .nl

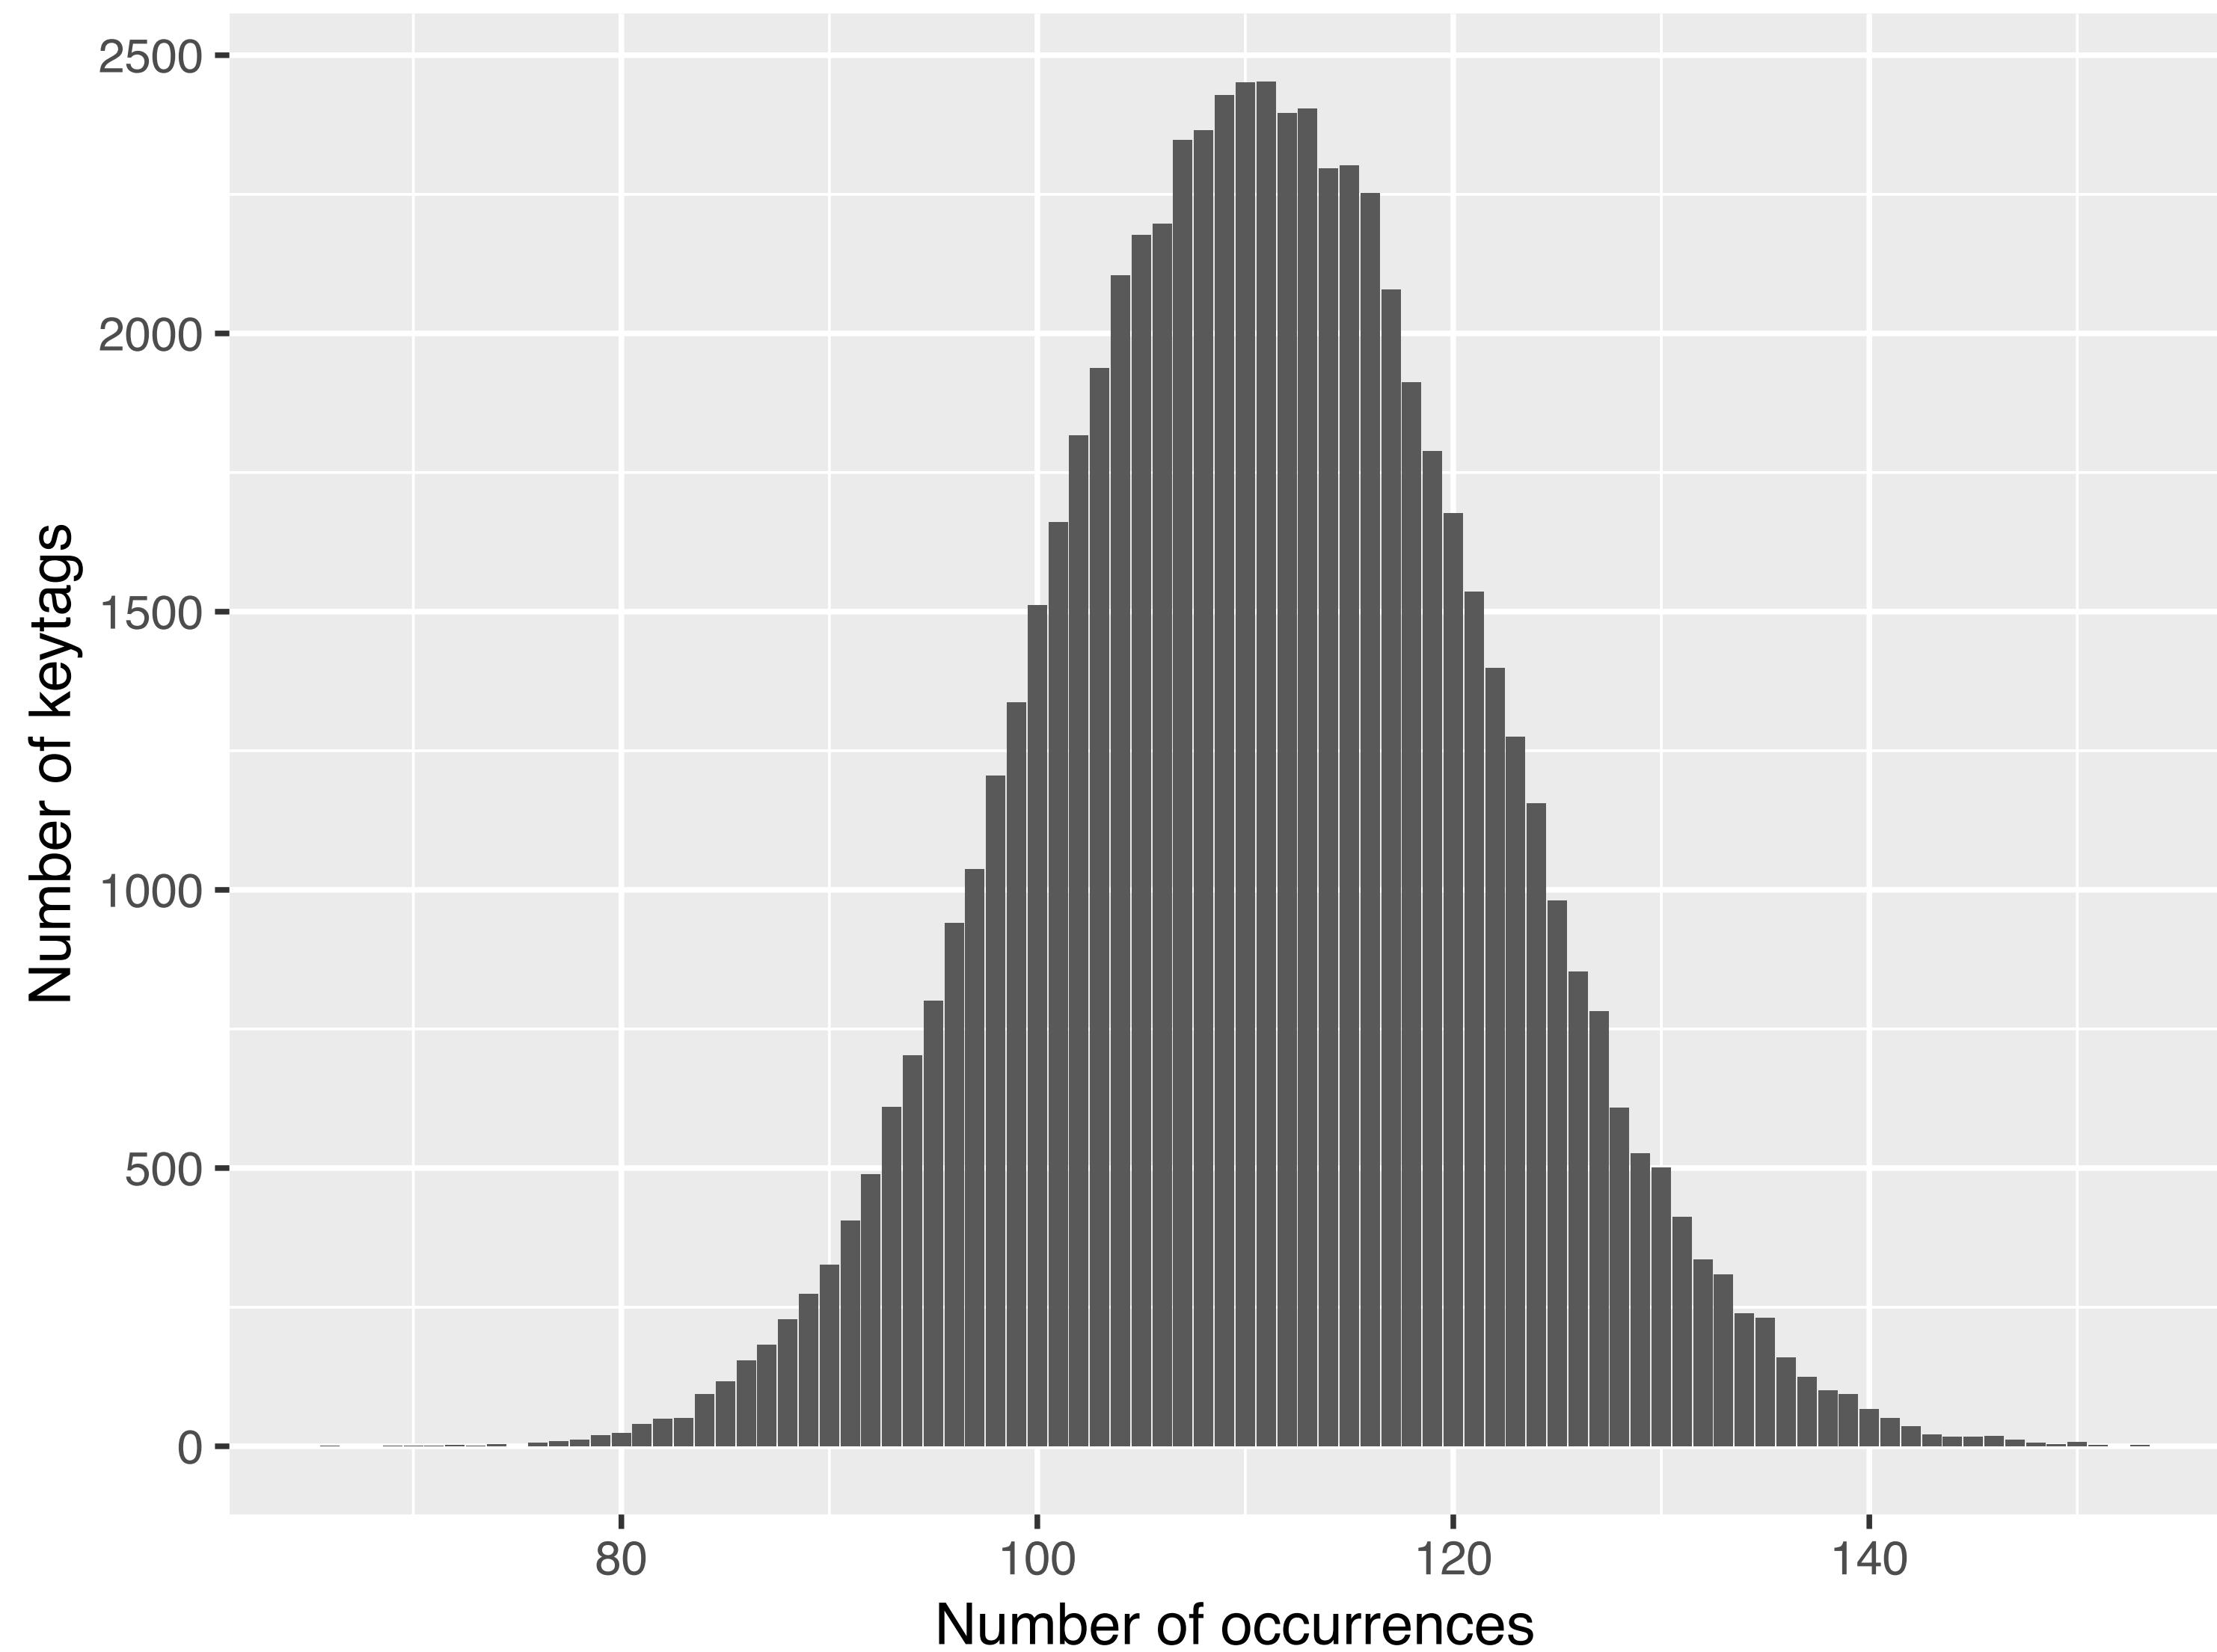


keytags

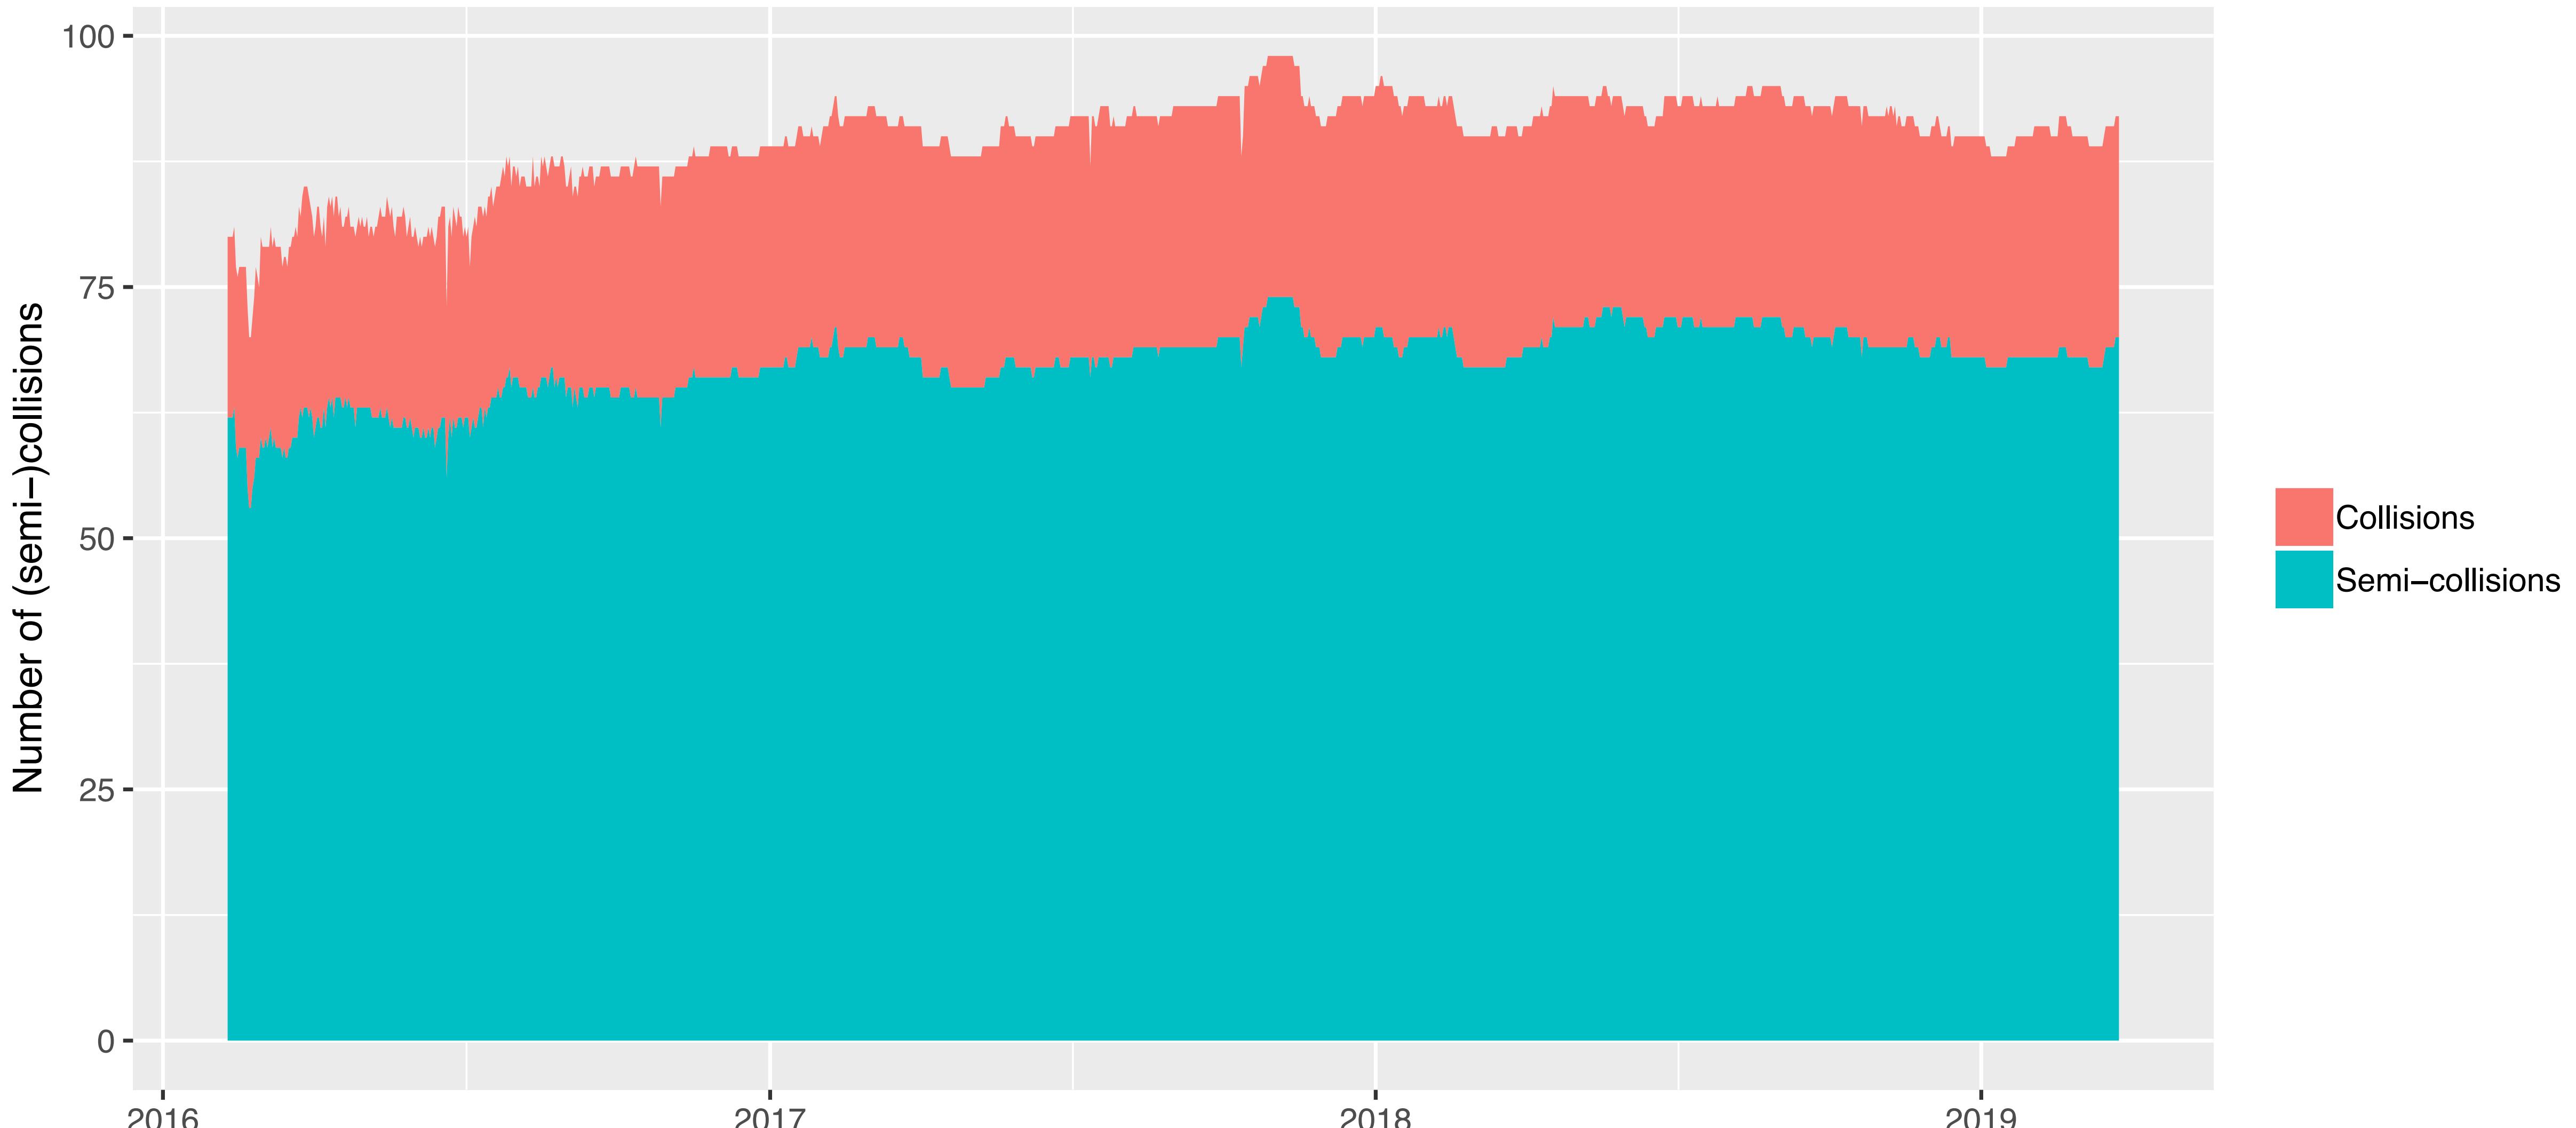


CRC16

# .nl RSA CRC16 distribution



# Is it actually better?



# Why is CRC16 not better?

- Random uniform (-like) distribution means; any keytag is equally likely
- Due to the Generalised Birthday Paradox, the likelihood of a collision increases with the number of keys in a keyset

#keys in DNSKEY set	Theoretical probability	Observed probability
2	0,00153%	0,00165%
3	0,00458%	0,00465%
4	0,00915%	0,02349%
<b>5 or more</b>	0,01526%	0,00853%

# Open questions

- The unspoken assumption for the empirical data is that implementations don't already prevent collisions from occurring
- For some, this assumption may not hold (e.g. LDNS, BIND)
- Can we fingerprint the crypto libraries used based on what we know from Roy's presentation and the distribution we observe?

# Conclusion: what did we learn?

- At first glance, the original keytag algorithm seems suboptimal
- Yet choosing something "better" in terms of the likelihood of values occurring turns out to be worse!
- This is at least an interesting lesson in protocol design; without really meaning to, the writers of the DNSSEC RFCs picked a "better" algorithm
- If we literally translate a Dutch phrase for this, they were:  
"Unknowingly capable" ;-)

# Operator and implementer advice

- Operators: if your domain is important and likely to be queried (very) frequently, then make sure you have no keytag collisions in your DNSKEY set -- the probabilities show that the chances of requiring more than one extra key generation is vanishingly small
- Implementers: consider (optionally) checking DNSKEY sets for keytag collisions and regenerate key(s) if a collision occurs
- Arguably: follow Postel's Law!



# Thank you! Questions?

 [nl.linkedin.com/in/rolandvanrijswijk](https://nl.linkedin.com/in/rolandvanrijswijk)

 [@reseauxsansfil](https://twitter.com/reseauxsansfil)

[roland@nlnetlabs.nl](mailto:roland@nlnetlabs.nl)